Repairing Safety After Benign Fine-Tuning: A Compute-Matched Comparison of Data Repairs

Anonymous Authors[1]

Abstract

Fine-tuning a safety-tuned language model on benign instruction data can erode its refusals of harmful requests. When only the training data may change, the owner must choose a repair: remove some rows, or add safety data. Many repairs have been proposed, but they have rarely been compared on equal terms, with training seeds, side effects and their compute counted. We measure which data repair restores the most safety for the compute spent without raising over-refusal or losing task skill. After benign fine-tuning of OLMo-2-1B-Instruct with a fixed LoRA recipe that erodes it partly, six pre-registered repairs, each run with six training seeds, are compared with random removal of the same size and with no repair: removal of rows chosen by training-data attribution, by a representation filter or by a free answer-length rule, and addition of random, selected or self-generated refusal examples. Each repair is labelled better, equivalent within 5 percentage points, worse, or inconclusive, and reported with its over-refusal, task loss and measured cost. Five references that change no data set the bar that a data repair must clear: a safety prompt, a down-scaled adapter, a post-hoc projection of the adapter, a gentler learning-rate schedule, and a published objective that constrains updates on the first answer tokens. A second family, Llama-3.2-1B-Instruct, is planned at a matched erosion dose. Pre-registered gates test whether the erosion is localized in data categories, a positive control tests whether diagnostics find a planted cause [PLACEHOLDER: if funded], and model-based predictions of when diagnosis pays are written down in advance and scored on the second family. We find [FINDING 1: the labels of the six repairs]; [FINDING 2: how the best data repair compares with the references]; and [FINDING 3: whether the erosion is localized, and whether diagnostics find the planted cause].

[1]Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

1. Introduction

Fine-tuning a safety-tuned ("Instruct") language model on data that contains nothing harmful can still erode its refusals of harmful requests. Qi et al. (2024) observed this after one epoch on ordinary instruction data sets (Alpaca, Dolly) for Llama-2-7b-Chat and GPT-3.5 Turbo; He et al. (2024) showed that benign rows selected for their similarity to harmful examples erode safety far more than random rows; and Guo et al. (2026) report that 100 random benign rows of full fine-tuning raise attack success sharply on recent 8B-class Instruct models. The size of the effect depends on the recipe (Kim et al., 2025). When a fine-tune does erode safety, its owner often controls only the data: commercial fine-tuning services screen training data, gate the tuned model on a safety evaluation, and report harm categories rather than rows, leaving the customer to fix the data and resubmit (Microsoft, 2026).

Repair is then a data decision with side effects. One can *remove* rows: published filters score benign rows by their contribution to erosion (He et al., 2024; Guan et al., 2025; Li et al., 2025; Wang et al., 2026b; Zhang et al., 2026b), and training-data attribution ranks rows by gradient similarity to harmful examples (Jiao et al., 2025; Pan et al., 2026). One can *add* safety data, but which data are added decides the outcome: on Llama3-8B-Instruct, averaged over seven downstream tasks, a random 1% mix of safety rows raised attack success from 20.24% (no safety data) to 59.23%, while DataRx's selected 1% mix lowered it to 13.70% (Zhang et al., 2026a). Safety data can also cause over-refusal (Bianchi et al., 2024; Luo et al., 2026b). And some repairs need no data change at all: a safety prompt used only at test time (Lyu et al., 2024), a gentler recipe (Kim et al., 2025), a post-hoc projection of the fine-tuned adapter (Hsu et al., 2024), or a fine-tuning objective that constrains updates on the first answer tokens (Qi et al., 2025). Every data repair costs compute, and the only direct evidence about it

is a retrain.

Which repair works best is an empirical question that published results do not settle. Each repair was evaluated in its own setting, with its own models, judges and baselines, mostly with one or a few training seeds, and seldom with its scoring compute charged against the alternative of retraining [CITE-CHECK]. The closest comparison we found, DualSelect (Chen et al., 2026), compares several selection methods with three seeds and a false-refusal rate [CITE-CHECK], but not addition repairs, model-level references or measured cost. Three results suggest that expensive diagnosis may not pay. Once selection compute is charged, powerful data-selection methods are almost never compute-optimal (Yin & Rush, 2025) [CITE-CHECK]. Gradient-based attribution tracks the answer format of training examples more than their task (Kim et al., 2026) [CITE-CHECK], so a free format or length rule may match it. And the erosion may be diffuse: even subsets selected for low conflict with safety still erode it (Guo et al., 2026), and Springer et al. (2026) argue from theory that better curation cannot remove the fragility.

We therefore measure repairs directly (§3, §4). On OLMo-2-1B-Instruct fine-tuned on 4,800 Dolly rows with a fixed LoRA recipe (recipe B: a constant learning rate, one epoch), six repairs fixed in advance are each run over six training seeds and compared with random removal of the same size and with no repair: three removals of 600 rows (ranked by training-data attribution, by a representation filter in the style of LARF, and by answer length) and three additions of 48 refusal examples (drawn at random, selected in the style of DataRx, or written by the untuned model itself). Each contrast is labelled *better*, *equivalent* (its 90% interval lies within ± 5 percentage points, pp), *worse* (significantly harmful on the unsafe rate, over-refusal or task loss) or *inconclusive*, after Holm correction. Over-refusal, measured against the untuned model, and task loss on held-out rows are feasibility limits. Cost is measured in GPU-seconds and in full-run equivalents (FRE) under two conventions. A data repair counts as dominant only if it beats five references that change no data, none of them ours: a safety prompt at test time, the fine-tuned adapter scaled down, its Safe LoRA projection (Hsu et al., 2024), recipe G (a recipe-only fix: warmup and cosine decay, which halve the mean learning rate), and recipe B with the constrained objective of Qi et al. (2025). A second family, Llama-3.2-1B-Instruct, is planned with recipe B at its own learning rate, chosen from erosion probes alone so that both families erode partly (a matched erosion dose; §4.1).

Three pre-registered parts surround the comparison. Gate A requires erosion beyond seed noise. Gate B asks whether the erosion is *localized*: does removing one Dolly category restore more safety than removing the same number of random rows? Gate C asks which diagnostic picks the culprit category, and whether a format rule explains attribution's pick. A positive control plants a known eroding cause in one cell of nested menus of 8, 32 and 128 cells [TO BE FINALISED] and asks whether each diagnostic finds it, and at which menu size diagnosis beats trial runs [PLACEHOLDER: if funded]. A secondary, model-based analysis predicts when paying for a diagnostic beats spending the same compute on trial runs selected by successive halving (Jamieson & Talwalkar, 2016; Li et al., 2018); its predictions are written down before Gate B's category runs and scored on the second family. A side study asks what the fine-tuning breaks: the refusal opening of an answer, the harmfulness of its content, or both. Every decision rule is committed before the data it judges; outcomes are replicated over training seeds, confirmed on held-out prompts, and reported as *too close to call* when the evidence cannot separate options.

Our contributions are:

- **A measured comparison of data repairs with equivalence bounds** (§3.2, §4.3–§4.4): six removal and addition repairs against random removal and no repair, over six training seeds each, labelled better, equivalent, worse or inconclusive, with over-refusal, task loss and measured cost, against five references that change no data, on OLMo-2-1B and, if its Gate A passes at a matched dose, on Llama-3.2-1B. [FINDING 1: the six labels, and which repair, if any, beats the references]. We release the menus, seeds, outcome tables and costs [ARTIFACT].

- **A positive control for diagnosis** (§3.4, §4.5): [FINDING 2: whether each diagnostic ranks a planted cause first, and the runs-only cost of finding it, at menu sizes 8, 32 and 128] [PLACEHOLDER: if funded].

- **Where the erosion lives** (§4.6, §4.10): [FINDING 3: the Gate B verdict with its detection limit; Gate C's reading with the format control; the share of the erosion that runs through refusal openings].

- **Pre-registered predictions of when diagnosis pays** (secondary; §3.5): [FINDING 4: predicted against measured value of diagnosis on the category and planted menus, and on the second family].

Section 2 positions the work. Section 3 defines the contrasts, labels, limits, the positive control and the secondary decision analysis. Section 4 gives the design. Section 5 reports results, including every too-close-to-call contrast. Section 6 discusses limitations, chiefly what the safety judge measures, the recipe and the scale.

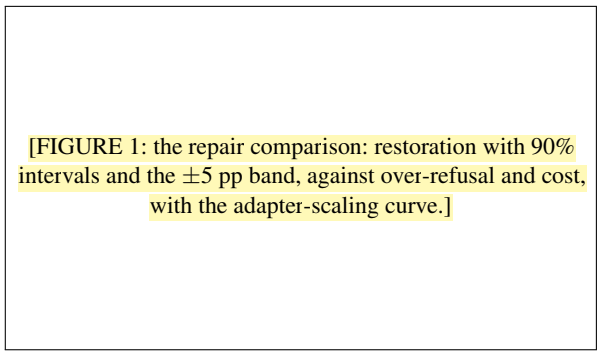


Figure 1. The repair comparison. For each repair, the safety restored against random removal or no repair, with its 90% interval and the ±5 pp equivalence band, plotted against over-refusal (left) and measured cost in FRE (right); the adapter-scaling reference curve is drawn in the left panel. (Placeholder: the figure is not yet drawn.)

2. Related Work

Safety erosion under benign fine-tuning. Qi et al. (2024) showed that fine-tuning aligned models on benign instruction data degrades safety, though less than harmful data does, with no training seeds reported. He et al. (2024) selected benign rows by representation and gradient similarity to harmful anchors; in their best case, 100 such rows gave over 70% unsafe answers against under 20% for random rows. This shows that benign rows *can* erode safety, not that ordinary fine-tunes have a localized cause. Kim et al. (2025) change the size of the erosion several-fold through optimization choices alone. Guo et al. (2026) find that 100 random benign rows of full fine-tuning erode recent 8B-class Instruct models, and that a handful of safety rows restores safety after a small attack. Under LoRA their results depend on the family: in their Table 2, Qwen2.5-7B-Instruct reads 7.27% before fine-tuning, 4.50% after 100 rows and 9.09% after 5,000 rows (two seeds), while Llama-3.1-8B-Instruct rises to 44.80% at 5,000 rows. Springer et al. (2026) explain erosion geometrically, as benign updates curving into a sharp, safety-relevant subspace.

Removal repairs: filters and attribution. Published filters score benign rows for their contribution to erosion. DataShield projects rows onto a compliance direction (Zhang et al., 2026b); LARF uses representations of safety-sensitive layers (Li et al., 2025); SQSD projects each row's parameter update onto danger and safety directions (Wang et al., 2026b); bi-anchoring compares a row's gradients or representations with harmful and safe anchors (He et al., 2024); and Self-Inf-N, proposed as an attack, finds outlier benign rows whose fine-tuning breaks safety (Guan et al., 2025) [CITE-CHECK]. Row filters report their scoring cost without charging it against runs (Zhang et al., 2026b; Hu et al., 2026) [CITE-CHECK]; other defences select tokens (Li et al., 2026) or use continual learning (Alssum et al., 2025) [CITE-CHECK]. DualSelect (Chen et al., 2026) selects task rows and safety references jointly and compares it with filters including LARF, TOSS and SQSD [CITE-CHECK]. For attribution, Pan et al. (2026) find that removing rows flagged by gradient-based attribution gives a safer retrained model than removing rows flagged by WildGuard, from one training run per arm and with no compute reported. DATE-LM (Jiao et al., 2025) finds that with safety-aligned look-alikes in the data, attribution's detection drops while LLM classifiers are largely unaffected; without them, the strongest classifiers detect better, yet removing attribution's picks gives the safer retrained model. Luo et al. (2026b) report weak, inconsistent effects from deleting or reweighting influential rows, large effects from rewriting them, and an over-refusal trade-off for safety rewriting.

Addition and model-level repairs. Adding about 3% safety examples to instruction tuning improves safety but can cause exaggerated refusals (Bianchi et al., 2024) [CITE-CHECK]. Which rows are added matters: DataRx selects safety rows by the gap between a model's own response and a safe reference (Zhang et al., 2026a) [CITE-CHECK], and its random mix made the model less safe than no safety data, while its selected mix made it safer (§1). Guo et al. (2026) restore safety with a handful of safety rows, and Llama 3's safety mix was tuned by repeated experiments that traded violations against false refusals (Grattafiori et al., 2024). Model-level repairs change no data: a safety prompt used only at test time reduces erosion (PTST; Lyu et al., 2024), SafeTune compares repairs of safety drift in single deterministic runs (Seth et al., 2026) [CITE-CHECK], a gentler recipe shrinks erosion (Kim et al., 2025), Safe LoRA projects the LoRA weights of selected layers onto a safety-aligned subspace without training or data (Hsu et al., 2024), and an objective that constrains updates on the first answer tokens makes alignment more persistent under fine-tuning (Qi et al., 2025). We use such repairs as references that a data repair must beat.

Attribution, its cost and what it tracks. Training-data attribution ranges from cheap gradient similarity on LoRA gradients (LESS; Xia et al., 2024) [CITE-CHECK] to estimates of finite removal effects by unrolled differentiation (SOURCE; Bae et al., 2024), which are reported to cost several training runs [CITE-CHECK]. Xiao & Aranguri (2026) compare the cost of diagnosis methods for undesirable post-training behaviours; probe attribution is over ten times cheaper than gradient attribution and LLM-judge baselines in their study. Yin & Rush (2025) charge selection compute against training compute and find that powerful selection methods are almost never compute-optimal; gradient-based selection pays only when the trained model is about ten times larger than the selection model [CITE-

CHECK]. Kim et al. (2026) find that gradient similarity tracks answer format more than task [CITE-CHECK]; we therefore include a free answer-length rule as a repair and as a control for attribution's picks.

Predicting outcomes and side effects. Persona vectors (Chen et al., 2025) project training data onto activation directions to flag data that will shift a trait [CITE-CHECK]. TuneAhead (Luo et al., 2026a) and Data2Behavior (Wang et al., 2026a) predict a fine-tune's outcome from its data before training, and MNEME predicts side effects of fine-tuning and unlearning by sparse model diffing (Kassem et al., 2025) [CITE-CHECK]. Any such predictor can enter our comparison as a scorer at its measured cost; side-effect prediction is a secondary analysis here (§4.10).

Decisions under a compute budget. Successive halving and Hyperband allocate a fixed budget across candidates by repeatedly discarding the worst at increasing fidelity (Jamieson & Talwalkar, 2016; Li et al., 2018); multi-fidelity methods model how cheap observations relate to full ones (e.g. Ouyang et al., 2025). DataDecide (Magnusson et al., 2025) scores small pretraining runs by decision accuracy against compute, and Ivison et al. (2025) charge selection compute against training on more data [CITE-CHECK].

How we differ. We do not claim a first benchmark of safety filters, a first comparison of removal against addition, or a first pricing of data selection: DualSelect, DataShield's cost table and Yin & Rush (2025) each cover part of this ground [CITE-CHECK]. Our contribution is the combination; to our knowledge, and not found by our search, no prior work combines equivalence-bounded contrasts over six training seeds for removal and addition repairs; over-refusal and task loss as limits; measured costs under two conventions; model-level references that a data repair must beat; a positive control with a planted cause; and predictions about the value of diagnosis written down before the data.

3. Problem Setup

3.1. Setting and endpoints

Let θ_0 be a safety-tuned Instruct model, fine-tuned with a fixed recipe $\mathcal{R}$ (adapter, loss, optimizer, schedule, length rule) on benign data $\mathcal{D} = D_1 \cup \cdots \cup D_K$, split into K *units* by metadata fixed before any outcome. A *repair* a maps $(\mathcal{D}, \mathcal{R})$ to a training set and recipe $(a(\mathcal{D}), \mathcal{R}_a)$; a data repair keeps $\mathcal{R}_a = \mathcal{R}$. The model trained with seed s (adapter initialization and data order) is $\theta(a, s)$. For a model θ, $f(\theta)$ is the *unsafe rate* (the share of sampled answers to a harm bank that a frozen judge labels unsafe), $\ell(\theta)$ is the held-out *task loss*, and $r(\theta)$ is the *over-refusal rate* on benign prompts that resemble harmful ones. Endpoints average over seeds: $F(a) = \mathbb{E}_s[f(\theta(a, s))]$, and likewise $\bar{\ell}(a)$ and $R(a)$. With a_0 the fine-tune without repair, the *erosion* is $E = F(a_0) - f(\theta_0)$; repair is meaningful only when E exceeds seed noise (Gate A, §4.6).

3.2. Repair contrasts, limits and labels (primary)

Let a_{R} remove the same number of random rows as the removal repairs. A removal repair m is compared with random removal, $\Lambda_m = F(a_{\mathrm{R}}) - F(m)$: the safety it restores beyond removing the same amount of data at random. An addition repair is compared with no repair, $\Lambda_m = F(a_0) - F(m)$, and a selected or self-generated addition with the random addition of the same size. The restoration $G_m = F(a_0) - F(m)$ is reported for every arm. Contrasts are estimated from per-seed unsafe rates as differences of means over unpaired seeds, with the variance pooled across arms.

Side effects are the *over-refusal excess* $O(m) = R(m) - r(\theta_0)$, against the untuned model ($R(m) - R(a_0)$ is also reported), and the *task-loss increase* $L(m) = \bar{\ell}(m) - \bar{\ell}(a_0)$. A repair is *feasible* at limits (τ_O, τ_L) if $O(m) \le \tau_O$ and $L(m) \le \tau_L$. Results are reported at $\tau_O \in \{2, 5, 10\}$ pp; τ_L is two standard errors of a seed-to-seed difference in held-out loss, fixed before the comparison runs. The *cost* of a repair is $C(m) = c(m) + 1$ FRE: $c(m)$ is the measured cost of building the repaired data (scoring rows, selecting or generating safety rows), and one FRE is the measured cost of one random-removal run (§4.7).

With a smallest effect of interest $\delta = 5$ pp and Holm correction at family-wise level 0.05 over the primary contrasts, each contrast receives one label:

- *better*: the Holm-adjusted test rejects $\Lambda_m = 0$ and the estimate is at least δ;
- *equivalent*: the 90% interval for Λ_m lies inside $(-\delta, \delta)$ (two one-sided tests);
- *worse*: the contrast with its comparator is significantly harmful at the Holm-adjusted 5% level on the unsafe rate ($\Lambda_m < 0$), on over-refusal ($R(m)$ above the comparator's) or on task loss ($\bar{\ell}(m)$ above the comparator's), each read in its harmful direction, so a harmful repair is never reported as inconclusive;
- *inconclusive*: otherwise.

Which label a repair receives when it is better on the unsafe rate but worse on a side effect is fixed in the stage-1 preregistration [TO BE FINALISED]; both results are reported. A margin of 4 pp is reported as a sensitivity. Five references change no data (§4.4). One, *adapter scaling*,

$$\theta_\lambda = \theta_0 + \lambda\,(\theta(a_0, s) - \theta_0), \qquad \lambda \in \{1, 0.75, 0.5, 0\},$$

traces (F, O, L) from the fine-tune ($\lambda = 1$) to the untuned model ($\lambda = 0$). A data repair *dominates* only if it lies below this curve: a lower F than the curve at equal O and L.

3.3. Localization

Let a_{-k} remove unit D_k and $a_{\mathrm{rand},k}$ remove $|D_k|$ random rows drawn across units. Erosion is *localized* at level δ_B if some $\Lambda_k = F(a_{\mathrm{rand},k}) - F(a_{-k})$ reaches δ_B beyond seed noise (the random arm controls for the smaller training dose), and *diffuse* if every Λ_k is credibly below δ_B. Diffuse erosion leaves the repair comparison unchanged: unit removals then act like random removal up to δ_B, and a unit-level diagnostic has little to find.

3.4. Positive control

A planted data set $\mathcal{D}^{\mathrm{p}}$ replaces the rows of one small cell of $\mathcal{D}$ with k rows that look benign but carry a known eroding signal, keeping the number of rows fixed. Nested partitions split $\mathcal{D}^{\mathrm{p}}$ into M cells, so that the planted cell sits inside one cell at every M. The planted erosion is $P = F(a_0^{\mathrm{p}}) - F(a_0)$, and removing the cell that holds the plant at menu size M restores $G^{\mathrm{p}}(M) = F(a_0^{\mathrm{p}}) - F(a_{-c(M)}^{\mathrm{p}})$. A diagnostic d *hits* at (M, k_d) if it ranks that cell among its top k_d. Because the truth is known, the sweep measures how diagnosis and trial runs scale with M, which natural menus at our budget cannot.

3.5. The value of diagnosis (secondary, model-based)

A *menu* $\mathcal{A}_M = \{a_0, a_1, \ldots, a_M\}$ holds a_0 and M repairs built by a frozen rule, never by a compared diagnostic; $\mathcal{A}_{\mathrm{feas}}$ is its feasible subset and $a^* = \arg\min_{a \in \mathcal{A}_{\mathrm{feas}}} F(a)$. A *selector* π with budget B (in FRE) buys trial runs (each a training fraction ϕ and a prompt fraction ψ, priced from the measured cost model) at total cost at most B and returns $\hat{a}$. It drops a repair whose adverse estimate of O or L passes its limit (a_0 is never dropped), and picks the lowest estimated F among repairs with full runs on at least two seeds; ties go to a_0. An infeasible pick is scored as if no repair were made:

$$\mathrm{loss}(\hat{a}) = \begin{cases} F(\hat{a}) & \hat{a} \in \mathcal{A}_{\mathrm{feas}}, \\ F(a_0) & \text{otherwise}, \end{cases}$$
$$\mathrm{Reg}(\pi; B, M) = \mathbb{E}[\mathrm{loss}(\hat{a})] - F(a^*).$$

This is a scoring convention for an expected outcome at a fixed budget, not a safety guarantee. The runs-only selector π_- uses one rule per (M, B) cell (a one-shot screen, successive halving or hierarchical bisection), chosen from a simulation before any outcome and then frozen. The selector π_{+d} pays the cost c_d of diagnostic d, keeps its top k_d repairs, and runs the same rule on $B - c_d$. The value of d and its break-even cost are

$$\Delta_d(B; M, c) = \mathrm{Reg}(\pi_-; B, M) - \mathrm{Reg}(\pi_{+d}; B, M, c),$$
$$c_d^*(B, M) = \sup\{c \geq 0 : \Delta_d(B; M, c) \geq 0\}.$$

$F(a^*)$ cancels in Δ_d. When a_0 is feasible, regret is non-negative, so no diagnostic can save more than $V_{\max}(B, M) = \mathrm{Reg}(\pi_-; B, M)$. At the menus we can afford, these quantities are mostly outputs of a model, so we treat them as predictions. A generative model (repair effects shrunk by empirical Bayes, noise from pooled residuals, measured costs, two effect shapes) is frozen before Gate B's category runs, and its predicted $V_{\max}$, Δ_d and c_d^* for the category menu and the planted menus are recorded then. We report measured against predicted values (replayed with the same code path on the outcome tables), and score OLMo-calibrated predictions of the chance of finding the best repair, and of regret, on Llama's category menu.

3.6. Opening decomposition (side study)

Let $p(\theta)$ be the share of sampled harm answers that open with a refusal (a frozen lexicon applied to the first 12 words), and $u_{\mathrm{R}}(\theta)$ and $u_{\mathrm{N}}(\theta)$ the unsafe rates among answers that do and do not open with one, so $f = p\,u_{\mathrm{R}} + (1-p)\,u_{\mathrm{N}}$. For the untuned model (subscript 0) and a fine-tune (subscript 1),

$$f_1 - f_0 = \underbrace{(p_0 - p_1)(u_{\mathrm{N}0} - u_{\mathrm{R}0})}_{\text{opening switch}} + \underbrace{(1 - p_1)(u_{\mathrm{N}1} - u_{\mathrm{N}0}) + p_1(u_{\mathrm{R}1} - u_{\mathrm{R}0})}_{\text{content change}}.$$

The *opening share* is the first term over $f_1 - f_0$; the split depends on the order of terms, so both orders are reported. The decomposition is descriptive: it says whether a change in the unsafe rate shows up in how answers open or in what they say, not which training choice caused the change. The first term has two layers on which families and recipes can differ: how many refusal openings fine-tuning removes $(p_0 - p_1)$, and how often a direct answer was already unsafe before fine-tuning ($u_{\mathrm{N}0}$); where $u_{\mathrm{N}0}$ is high and $u_{\mathrm{R}0}$ low, a lost opening comes close to an unsafe answer. If the opening share is small, repairs that act only on how answers open (refusal examples; an objective that constrains updates on the first answer tokens) may restore less than repairs that change what the answers say.

4. Experimental Design

4.1. Models, data and recipes

Families and scale. OLMo-2-0425-1B-Instruct (Team OLMo et al., 2024) is the stage-1 model; Llama-3.2-1B-Instruct (Grattafiori et al., 2024) is the planned second family, at a matched erosion dose (below; §4.8). Larger sizes

follow a ladder: a published-setting replication and a reduced study at 7–8B, then a headline check at 32B or 70B [PLACEHOLDER: sizes actually run]. Qwen2.5-1.5B- and 7B-Instruct (Qwen et al., 2024) are outside the repair study; their existing recipe-B runs enter only the same-recipe family contrast (§4.10). Every model trains on one GPU, and every comparison stays on one GPU type.

Data. Gate B and stage 1 train on 4,800 rows of Dolly-15k (Conover et al., 2023), 600 from each of 8 task categories (the units; $K = 8$), prepared and sealed per tokenizer with a manifest; 200 further rows, 25 per category, measure task loss. The add repairs draw their safety rows from PKU-SafeRLHF (§4.3). Alpaca is the planned second benign dataset, after Gate B [NOT YET FIXED: its units and learning rate].

Recipes. *Recipe B*, the main recipe: LoRA (Hu et al., 2022) of rank 64 and scale 128 on all linear layers; AdamW at a constant learning rate of 10^{-4}, no warmup; 128 rows per step; one epoch (38 steps); a per-row loss (each row's mean over its response tokens). Gate A, Gate B and stage 1 use it. It was chosen over recipe G from probe and Gate A readings alone, before any repair outcome. *Recipe G*, the recipe-only fix, adds a 4-step warmup and cosine decay, which halve the mean learning rate over the epoch, and a token-mean loss; it is a reference (§4.4). A pre-registered single-seed run that changes only the loss weighting separates its two changes (§4.10). Training is deterministic on matched hardware; a seed sets the adapter initialization and the data order.

A matched erosion dose. Each family uses recipe B with only its learning rate changed: the highest rung of a ladder (10^{-4}, 5×10^{-5}, 2×10^{-5}) whose erosion probe reads partial (20–55 pp) under the frozen one-seed versions of the Gate A bars (§4.6; a too-close probe gets one more seed). The rate is chosen from probes only, never from a repair result. During development no rung of a one-recipe ladder left both families partly eroded (§5), because families respond to the same change of recipe by different amounts; we fixed this rule after reading that ladder and the second family's first probe at recipe B, and state cross-family claims at a matched erosion dose. OLMo-2-1B's rung is 10^{-4}. [PLACEHOLDER: Llama-3.2-1B's rung.]

4.2. Measures, judges and the claim rule

- **Unsafe rate U.** The share of 600 sampled answers (200 development harm prompts $\times$ 3 samples at temperature 0.7, up to 2,048 new tokens) that Qwen3Guard-Gen-4B (Zhao et al., 2025) labels Unsafe. The prompts come from HarmBench (Mazeika et al., 2024), StrongREJECT (Souly et al., 2024), SimpleSafetyTests (Vidgen et al., 2023) and OR-Bench's toxic set (Cui et al., 2025). Standard errors are clustered by prompt. U is a broad policy label, not an attack-success rate as published elsewhere: judges and prompts differ, so U is compared only within this study.

- **Compliance and the claim rule.** HarmBench's classifier asks whether an answer carries out the requested behaviour. A contrast in U supports a sentence about *harm* only if the classifier's contrast has the same sign and at least half the size; otherwise the sentence says *restored refusal*. If the classifier's erosion under a recipe is below 20 pp, no sentence calls that recipe's erosion harm.

- **Also per arm:** hard-category U (violent, illegal, sexual or self-harm categories); U by answer-length band and a length-reweighted U; Granite Guardian 3.3 (Padhi et al., 2024) on the greedy answers; and a selectivity score (refusal on harm prompts minus refusal on benign prompts).

- **Over-refusal.** Qwen3Guard's refusal flag on greedy answers to 100 benign prompts that resemble harmful ones (50 XSTest safe prompts, Röttger et al. 2024; 50 OR-Bench-Hard prompts, Cui et al. 2025), reported by source, with a lexicon check.

- **Task loss and skill.** Held-out loss on all 200 task rows is the skill limit of the comparison. The probe and Gate A bars also check ARC-Easy accuracy (Clark et al., 2018) on 189 questions: the paired drop d against the untuned model must stay at or below 0.03. With 189 questions the paired interval is about ± 0.07 wide, so it straddles 0.03 for almost every run. For probe and Gate A readings, a rule fixed during development, before the readings it was written for, resolves a straddling interval [TO BE FINALISED: its scope under the per-family rule]: the run **passes, as "skill kept, reported as a limitation", when the point drop $d \leq 0.03$ and the held-out task loss is below the untuned model's (the task is learned); it fails when $d > 0.03$.** This rule is a loosening: it can pass a run whose true drop is up to about 6 points, which 189 questions cannot detect. We report the verdict under the frozen bars beside the verdict under this rule. Disclosure: the first recipe-G probe on OLMo-2 was read under the earlier rule before this rule was proposed, so for that one run the rule is not applied blind. At 1B, ARC-Easy is close to chance, so "skill kept" rests mainly on the held-out task loss.

- **Answers end.** A coherence screen flags answers that reach the token cap, repeat 4-grams heavily, are empty, or have fewer than 5 tokens; an arm whose answers fail too often is not read.

- **Banks.** Gates and selectors read development banks only. Final claims are confirmed on a sealed held-out subsample (§4.8).

4.3. Stage 1: the primary comparison

Table 1 lists the stage-1 arms. Every rule, scorer and pool is committed by hash before Gate B's category runs (§4.9).

Primary family. P1: A1 against a_{R}; P2: L1 against a_{R}; P3: H1 against a_{R}; P4: S1 against a_0 (two-sided, since prior work predicts both directions; a significant harmful result is labelled *worse*); P5: S3 against S1; P6: S4 against S1. Holm at 5%, labels as in §3.2.

Pools and anchors. The refusal pool holds prompts from PKU-SafeRLHF (Ji et al., 2025) that have an unsafe response in the source data, paired with a response that the source labels safe and that opens with a refusal by the frozen lexicon. Pools and anchors are never filtered by the evaluation judge, and prompts that nearly duplicate an evaluation prompt are removed. The add effect must also hold on the evaluation prompts least similar to the added rows.

Seeds and power. Method arms run at six seeds, random removal at eight and no repair at six; seeds run in blocks fixed before any outcome, and only complete blocks are analysed. A blinded rule reads only Gate B's pooled noise estimate: up to 4.5 pp, the plan above; from 4.5 to 6.0 pp, the family shrinks to P1, P3 and P6 at nine seeds; above 6.0 pp, stage 1 does not run. At planning noise the minimum detectable effect after Holm's first step is about 5 pp [PLACEHOLDER: final MDE from Gate B's pooled noise].

Secondary and scored-only arms [TO BE FINALISED]. At three seeds each, exploratory, with one-sided bounds only: a random dose curve (1,200 and 2,400 rows removed); random removal plus S1; a DataShield-style filter; rules for answers with list or markup lines and for the longest answers; the length rule under a token-mean loss; and octiles of answer length. Scored without new runs: a Self-Inf-N-style score and a content classifier; we publish the overlap of all top-600 sets. SQSD is cited but not run, because its scoring needs several FRE.

4.4. References that change no data

All five are mandatory, and we claim no novelty for any. **Safety prompt:** a fixed safety instruction as a prefix of the user turn (OLMo-2 has no system turn), at evaluation only, as in PTST (Lyu et al., 2024), on three a_0 seeds. **Adapter scaling:** the a_0 adapter scaled by $\lambda = 0.75$ and 0.5 on three seeds, which with $\lambda \in \{0, 1\}$ gives the reference curve of §3.2. **Safe LoRA** (Hsu et al., 2024): a_0's adapter projected as in §2, on three seeds; the base model, layer rule and threshold follow the paper's rule or a fixed default, never stage-1 outcomes [NOT YET FIXED]. **Recipe G** (recipe only): the unchanged data under recipe G, three seeds. **B+Qi:** recipe B with the token-wise constrained SFT objective of Qi et al. (2025), which constrains updates on the first answer tokens, on the unchanged data at six seeds. The implementations of Safe LoRA and B+Qi are committed before stage 1 is frozen. A pinned evaluation script regenerates each a_0 adapter by a deterministic re-run and is certified once: with $\lambda = 1$ and no prefix it must reproduce a_0's U exactly. The paper states plainly whether any data repair beats these references.

4.5. Positive control: a planted cause

[PLACEHOLDER: runs only if funded; otherwise menu size leaves the headline.] The plant is $k = 37$ rows whose answers comply directly with prompts that also have an unsafe answer in the source data, but which the source labels safe; they look benign. A sealed draw picks one category and one leaf cell inside it, and the plant replaces that cell's rows, so the run keeps 4,800 rows and 38 steps. The nested cells [TO BE FINALISED] are the 8 categories (600 rows), 32 length quartiles within categories (150 rows) and 128 random quarters of those (37–38 rows). A pilot of three seeds goes on only if the plant adds at least 10 pp of erosion with a one-sided 90% lower bound above 5 pp (otherwise the plant doubles once, then the null is reported). The main sweep removes the planted cell at each menu size (three seeds each) and runs random removal on the planted data. Every scorer ranks the cells without new runs. Outputs: $G^{\mathrm{p}}(M)$; each diagnostic's top-k hit rate with known truth; the runs-only cost of finding the cell, replayed with the frozen selector rules; and a check of additivity across the nested levels.

4.6. Gates A, B and C

Gate A (entry). For each family at its matched dose, over three seeds (the probe seed and two fresh seeds), mean E must lie in a pre-registered partial-erosion window of 20–55 pp, be at least twice the seed standard deviation, and exceed zero on every seed; a second judge must agree in direction, the answers must end, and the task must be learned without lost skill (§4.2). [PLACEHOLDER: Gate A result with every setting tried and its disclosures (§5).]

Gate B (localization). These rules summarize its pre-registration, revision 5 [NOT YET FIXED]. At recipe B, each category removal (600 rows, 33 steps) is compared with random removal of 600 rows (75 per category) at the same seed, so a_0 cancels. All 8 categories run at one seed; the two with the largest contrast, plus Gate C's picks, run at two more seeds. The verdict is LOCALIZED if some category restores at least $\delta_B = 8$ pp more than random

Table 1. Stage-1 arms on OLMo-2-1B (recipe B; one epoch; removal arms train 33 steps, addition arms 38). The references of §4.4 are listed below the repairs.

Arm	Repair	Rule	Seeds
a_{R}	random removal	600 random rows, 75 per category	8
a_0	no repair	–	6
H1	length rule (free)	remove the 600 rows with the shortest answers	6
A1	attribution	remove the top 600 rows by bi-anchored gradient similarity at two checkpoints of a_0 (the Gate C method)	6
L1	LARF-style filter	remove the top 600 rows by the untuned model's representation similarity to harmful anchors at one layer	6
S1	random refusals	add 48 rows (1%) drawn at random from a fixed refusal pool	6
S3	DataRx-style	add 48 rows selected from S1's pool by DataRx's score	6
S4	self-generated	add 48 pool prompts with the untuned model's own refusals	6
B+Qi	constrained SFT (reference)	recipe B with the token-wise constrained SFT objective of Qi et al. (2025)	6
G	recipe only (reference)	the unchanged data, trained with recipe G	3
Safe LoRA	projection (reference)	project a_0's adapter as in Hsu et al. (2024) [NOT YET FIXED: settings]	3
prefix	safety prompt (reference)	evaluate a_0 with a safety instruction prefixed to the user turn (no training)	3
λ	adapter scaling (reference)	evaluate a_0's adapter scaled by $\lambda = 0.75$ and 0.5 (no training)	3

removal and clears a Bonferroni-corrected bar; DIFFUSE if every category is credibly below 8 pp; and TOO CLOSE otherwise. Guards require answers that end, recompute the verdict on coherent answers and under adverse counting of unscored answers, and reject a restoration that over-refuses. A noise gate stops the stage if seed noise is too large; an automatic extension adds two seeds for contested categories; and a length-matched control separates content from answer length for a localized category. By simulation at planning noise, an 8 pp category reads LOCALIZED about half the time and 80% detection needs about 10–12 pp; DIFFUSE therefore means no category-level cause above about 8 pp, not no cause.

Gate C (which diagnostic picks the category). On the first seed's a_0, attribution (A1's score) and representation similarity each pick their top category; a format rule picks the category with the largest share of its loss on the first five answer tokens. Gate C is read only if Gate B is LOCALIZED (AGREE, ATTR_BETTER, REP_BETTER, NONE_FINDS or TOO CLOSE). If attribution's pick equals the format pick, the paper says that answer form explains it.

If DIFFUSE. Gate C does not run and stage 1 is unchanged. A rewrite arm (the untuned model's own answers replace the 600 top-attribution rows) runs first among the secondary arms, and the planted sweep becomes the only test of whether diagnosis can pay.

4.7. Cost accounting

Every arm and every scorer is logged in GPU-seconds by stage (setup, training, greedy generation, sampled generation, judging, scoring, overhead). One FRE per family is the median measured time of that family's random-removal arms. Two cost conventions are frozen before stage 1: *full* (training plus the evaluation as run) and *lean* (training plus one sample per harm prompt, priced from the measured time per answer); the cost c_0 of one lean trial run is measured on a real job. Every pool filter, generated answer, anchor job and scorer is charged; "free" means below 0.01 FRE. Attribution is priced from its measured scoring pass. Claims about cost hold only where both conventions agree, and no cost ratio measured at 1B is claimed for a larger model. [TO BE FINALISED: measured costs per arm and per scorer.]

4.8. Second family, confirmation and scale

Llama-3.2-1B (stage 2) is planned at its matched dose (§4.1): its own Gate A, a reduced Gate B on categories, and the primary family at six seeds. If both families run, a contrast is said to "replicate across two families" only if it has the same sign in both and the family difference is below the interaction's MDE; a family difference is claimed only above that MDE. The same-recipe family contrast (§4.10) is not a replication. **Confirmation** (stage 3): at most two contrasts per family, chosen by a rule committed before stage 1, are re-run at six fresh seeds and evaluated on a sealed 200-prompt held-out subsample as well as the development bank. **Scale** (stage 4): first a replication in a published setting (a LARF- or DataShield-style filter on Llama-3.1-8B-Instruct fine-tuned on Alpaca, with random removal and length rules), then the scale ladder [PLACEHOLDER: if resources allow]. No result is claimed at a size without runs at that size.

4.9. Pre-registration and testing order

Gate B runs first, to its verdict; funding for the later stages is then sought with that result, and stage 1 starts only once funded, never because of an outcome. Before Gate B's category runs, the stage-1 pre-registration [NOT YET FIXED]

commits every stage-1 scorer's code and settings hash, every edit rule, the pools, the plant, the primary family, the seed rule, the limits, the claim rule, the selector specification and its predictions. Anything added later is labelled post-Gate-B exploratory. Families of tests are read in order, each at its own 5% level: (1) the primary family; (2) confirmation; (3) the second-family (Llama) comparison; (4) the planted-sweep claims; (5) the decision-value predictions. Everything else is descriptive.

4.10. Side studies

1. **Mechanism: refusal openings against harmful content.** The opening decomposition (§3.6) on every arm and recipe, with a frozen refusal lexicon, reports how much of each change in U comes from answers that stop opening with a refusal and how much from what the answers say. It locates a change; it does not say which training choice caused it. Across families and recipes we report both layers of §3.6: the refusal openings lost, and the unsafe rate of direct answers before fine-tuning. One causal test runs on Gate B arms: a teacher-forced first-token refusal margin $m(p) = \log P_{\text{ref}}(p) - \log(1 - P_{\text{ref}}(p))$, where P_{ref} is the probability of a frozen refusal phrase as the first answer tokens; Qi et al. (2025) argue that safety alignment often adapts mainly the first few output tokens. If a category's effect runs through openings, the margin should track U across arms. [PLACEHOLDER: report against the bound on a first-token path recorded before Gate B.] Exploratory.

2. **Recipe and family contrasts.** Recipe G against recipe B measures a recipe-only fix (§4.4); a pre-registered single-seed run with recipe B's schedule and recipe G's token-mean loss separates G's two changes [PLACEHOLDER: its result against its pre-registered prediction, §5]. We also report the same-recipe family contrast (recipe B on OLMo-2, Llama-3.2 and Qwen2.5), every learning-rate rung that ran, and a pre-registered rank-16 check under recipe G. This is context, not a main claim.

3. **Measurement validity.** Three judges on the same answers (Qwen3Guard, Granite Guardian, HarmBench's classifier), with agreement, E per judge, hard-category U and length bands, and an account of how much of the gap to published numbers comes from judges and prompts rather than training.

4. **Over-refusal as the cost of repair.** O against the untuned model for every arm, by source; the judge's own false-Unsafe rate on benign prompts bounds its precision.

5. **Side-effect prediction (secondary).** On the addition arms only: the balanced error of predicting an over-refusal violation, a diagnostic against trial runs of equal cost.

5. Results

[PLACEHOLDER: Gate A (entry), with every setting tried and the audit's disclosures; Table 2, the stage-1 comparison (G, Λ with 90% intervals and labels, O, L, cost under both conventions, the HarmBench contrast and the claim-rule wording), with the five references and the adapter-scaling curve; Figure 1; the second-family comparison at its matched dose, if run; the same-recipe family contrast, the development ladder and every learning-rate rung, the token-mean-only run against its pre-registered prediction and the rank-16 check; Gate B's verdict with its detection limit and the length control; Gate C and the format reading; the planted sweep, if funded; measured against predicted decision value; the opening decomposition with both layers, and the first-token margin; every too-close-to-call result.]

6. Discussion and Limitations

What "unsafe" means. Our primary judge gives a broad policy label. After fine-tuning it may track whether a model refused more than whether its answer carries out harm. The claim rule (§4.2) therefore decides between "restored refusal" and "reduced harm" for each contrast, and U is not comparable with attack-success rates published under other judges and prompts.

Equivalence and power. The labels resolve differences of about 5 pp; smaller differences stay inconclusive, and an *equivalent* label means "within ± 5 pp", not "identical". Gate B's detection limit (§4.6) holds for categories; row-level causes that span categories are not tested by Gate B.

Recipe and setting. The main recipe, B, is plain rather than standard: a constant learning rate with no warmup and a per-row loss, with a rank-64 adapter on all linear layers, larger than is common (a rank-16 check is reported). A standard warmup and cosine schedule (recipe G) lowers the mean learning rate; we report it as a recipe-only fix, and repair rankings may differ at a lower dose. A second family runs at its own learning rate, so cross-family claims hold at a matched erosion dose, not at one recipe; that rule was fixed after a development ladder was read. The skill bar at the probes and Gate A is looser than the bar first registered, and we say where it decided a reading (§4.2). The main model and recipe were chosen by screening probe and Gate A readings, never Gate B or stage-1 outcomes; §5 lists every setting tried.

Baselines and doses. Every scorer is re-implemented from its paper ("X-style"), so a weak result may reflect our implementation; a published table of deviations per scorer, reference outputs requested from the authors and the published-setting replication address this. Comparisons across removal and addition hold only at the pre-registered doses (600 rows removed, 48 added).

Cost. An arm's cost depends on how much we evaluate as well as on training. We report two cost conventions and claim only where they agree. Costs at 1B do not transfer to larger models, because a scorer's cost and a run's cost grow differently with model size.

Scale and family. The full comparison runs at 1B on OLMo-2, and on Llama-3.2 if its Gate A passes at its matched dose, with one benign dataset until Alpaca runs; larger sizes get a replication only if resources allow. A family difference is claimed only above its MDE.

The positive control and the secondary analysis. The planted cause is artificial; finding it shows that a diagnostic *can* find a concentrated cause, not that natural erosion has one. The value of diagnosis at natural menus is model-based and judged against predictions made in advance.

Scope and user. We compare data repairs with model-level references; rewriting responses enters only if the erosion is diffuse. A customer of a closed fine-tuning service can change only the data but cannot run gradient or representation scorers; a developer with open weights can run every scorer but also has model-level repairs. Our answer serves a service that must ship the customer's adapter and may change only the data, and a pipeline that allows only data changes.

7. Conclusion

[PLACEHOLDER: which repair to use, at what cost and with what side effects, and whether paying for diagnosis was ever worth it, under the observed outcome.]

References

Alssum, L., Itani, H., Hammoud, H. A. A. K., Torr, P., Bibi, A., and Ghanem, B. Unforgotten safety: Preserving safety alignment of large language models with continual learning. arXiv preprint arXiv:2512.10150, 2025. URL https://arxiv.org/abs/2512.10150.

Bae, J., Lin, W., Lorraine, J., and Grosse, R. Training data attribution via approximate unrolled differentiation. arXiv preprint arXiv:2405.12186, 2024. URL https://arxiv.org/abs/2405.12186.

Bianchi, F., Suzgun, M., Attanasio, G., Rottger, P., Jurafsky, D., Hashimoto, T., and Zou, J. Safety-tuned LLaMAs: Lessons from improving the safety of large language models that follow instructions. In *The Twelfth International Conference on Learning Representations*, 2024. URL https://openreview.net/forum?id=gT5hALch9z.

Chen, R., Arditi, A., Sleight, H., Evans, O., and Lindsey, J. Persona vectors: Monitoring and controlling character traits in language models. arXiv preprint arXiv:2507.21509, 2025. URL https://arxiv.org/abs/2507.21509.

Chen, X., Zhang, J., Wu, O., and Gao, D. Two to tango: Coupled task-reference selection for safe LLM fine-tuning. arXiv preprint arXiv:2606.09866, 2026. URL https://arxiv.org/abs/2606.09866.

Clark, P., Cowhey, I., Etzioni, O., Khot, T., Sabharwal, A., Schoenick, C., and Tafjord, O. Think you have solved question answering? Try ARC, the AI2 reasoning challenge. arXiv preprint arXiv:1803.05457, 2018. URL https://arxiv.org/abs/1803.05457.

Conover, M., Hayes, M., Mathur, A., Xie, J., Wan, J., Shah, S., Ghodsi, A., Wendell, P., Zaharia, M., and Xin, R. Free Dolly: Introducing the world's first truly open instruction-tuned LLM. Databricks blog, 2023. URL https://www.databricks.com/blog/2023/04/12/dolly-first-open-commercially-viable-instruction-tuned-llm.

Cui, J., Chiang, W.-L., Stoica, I., and Hsieh, C.-J. OR-Bench: An over-refusal benchmark for large language models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 11515–11542. PMLR, 2025. URL https://proceedings.mlr.press/v267/cui25a.html.

Grattafiori, A., Dubey, A., Jauhri, A., et al. The Llama 3 herd of models. arXiv preprint arXiv:2407.21783, 2024. URL https://arxiv.org/abs/2407.21783.

Guan, Z., Hu, M., Zhu, R., Li, S., and Vullikanti, A. Benign samples matter! Fine-tuning on outlier benign samples severely breaks safety. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 20572–20597. PMLR, 2025. URL https://proceedings.mlr.press/v267/guan25c.html.

Guo, Y., Chen, X., Zhang, S., Wang, X., and Tang, H. When safety routing breaks: Understanding alignment fragility under benign fine-tuning. arXiv preprint arXiv:2609.01455, 2026. URL https://arxiv.or

g/abs/2609.01455. Accepted to the Findings of EMNLP 2026.

He, L., Xia, M., and Henderson, P. What is in your safe data? Identifying benign data that breaks safety. In *First Conference on Language Modeling*, 2024. URL https://openreview.net/forum?id=Hi8jKh4HE9.

Hsu, C.-Y., Tsai, Y.-L., Lin, C.-H., Chen, P.-Y., Yu, C.-M., and Huang, C.-Y. Safe LoRA: The silver lining of reducing safety risks when finetuning large language models. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024. URL https://openreview.net/forum?id=HcifdQZFZV.

Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., and Chen, W. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022. URL https://openreview.net/forum?id=nZeVKeeFYf9.

Hu, Z., Huang, X., Mendoza, P., Alghamdi, E. A., Alomair, B., Popa, R. A., and Wagner, D. GradShield: Alignment preserving finetuning. arXiv preprint arXiv:2605.14194, 2026. URL https://arxiv.org/abs/2605.14194.

Ivison, H., Zhang, M., Brahman, F., Koh, P. W., and Dasigi, P. Large-scale data selection for instruction tuning. arXiv preprint arXiv:2503.01807, 2025. URL https://arxiv.org/abs/2503.01807.

Jamieson, K. and Talwalkar, A. Non-stochastic best arm identification and hyperparameter optimization. In *Proceedings of the 19th International Conference on Artificial Intelligence and Statistics*, volume 51 of *Proceedings of Machine Learning Research*, pp. 240–248. PMLR, 2016. URL https://proceedings.mlr.press/v51/jamieson16.html.

Ji, J., Hong, D., Zhang, B., Chen, B., Dai, J., Zheng, B., Qiu, T. A., Zhou, J., Wang, K., Li, B., Han, S., Guo, Y., and Yang, Y. PKU-SafeRLHF: Towards multi-level safety alignment for LLMs with human preference. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 31983–32016. Association for Computational Linguistics, 2025. URL https://aclanthology.org/2025.acl-long.1544/.

Jiao, C., Pan, Y., Xiao, E., Sheng, D., Jain, N., Zhao, H., Dasgupta, I., Ma, J. W., and Xiong, C. DATE-LM: Benchmarking data attribution evaluation for large language models. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2025. URL https://openreview.net/forum?id=e2cD5xuHix.

Kassem, A. M., Shi, Z., Rostamzadeh, N., and Farnadi, G. REVIVING YOUR MNEME: Predicting the side effects of LLM unlearning and fine-tuning via sparse model diffing. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 32250–32263. Association for Computational Linguistics, 2025. URL https://aclanthology.org/2025.emnlp-main.1641/.

Kim, M., Kwak, J. M., Alssum, L., Ghanem, B., Torr, P., Krueger, D., Barez, F., and Bibi, A. Rethinking safety in LLM fine-tuning: An optimization perspective. In *Second Conference on Language Modeling*, 2025. URL https://openreview.net/forum?id=ZnOoEA2nDn.

Kim, S., Jung, S., Kim, S., Oh, S. J., and Oh, A. Form over content in gradient-based data attribution methods. arXiv preprint arXiv:2609.19589, 2026. URL https://arxiv.org/abs/2609.19589.

Li, H., Li, L., Lu, Z., Wei, X., Li, R., Shao, J., and Sha, L. Layer-aware representation filtering: Purifying finetuning data to preserve LLM safety alignment. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 8030–8050. Association for Computational Linguistics, 2025. URL https://aclanthology.org/2025.emnlp-main.406/.

Li, L., Jamieson, K., DeSalvo, G., Rostamizadeh, A., and Talwalkar, A. Hyperband: A novel bandit-based approach to hyperparameter optimization. *Journal of Machine Learning Research*, 18(185):1–52, 2018. URL http://jmlr.org/papers/v18/16-558.html.

Li, Y., Liu, Z., Li, Z., Lin, Z., and Zhang, J. Token-level data selection for safe LLM fine-tuning. In *The Fourteenth International Conference on Learning Representations*, 2026. URL https://openreview.net/forum?id=k7ytptAaDN.

Luo, Y., Long, H., Wang, C., Duan, Q., Lin, X., Xu, Y., Luo, Y., Yang, W., and Tang, N. TuneAhead: Predicting fine-tuning performance before full training begins. In *Forty-third International Conference on Machine Learning*, 2026a. URL https://openreview.net/forum?id=JBFB2pUOCA.

Luo, Y., Wang, C., Chen, J., and Pan, L. From reweighting to rewriting: Unlocking the intervention effects of influential samples in training data attribution. arXiv preprint arXiv:2609.02771, 2026b. URL https://arxiv.org/abs/2609.02771.

Lyu, K., Zhao, H., Gu, X., Yu, D., Goyal, A., and Arora, S. Keeping LLMs aligned after fine-tuning: The crucial role of prompt templates. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*,

2024. URL https://openreview.net/forum?id=xNlQjS0dtO.

Magnusson, I., Tai, N., Bogin, B., Heineman, D., Hwang, J. D., Soldaini, L., Bhagia, A., Liu, J., Groeneveld, D., Tafjord, O., Smith, N. A., Koh, P. W., and Dodge, J. DataDecide: How to predict best pretraining data with small experiments. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 42487–42502. PMLR, 2025. URL https://proceedings.mlr.press/v267/magnusson25a.html.

Mazeika, M., Phan, L., Yin, X., Zou, A., Wang, Z., Mu, N., Sakhaee, E., Li, N., Basart, S., Li, B., Forsyth, D., and Hendrycks, D. HarmBench: A standardized evaluation framework for automated red teaming and robust refusal. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 35181–35224. PMLR, 2024. URL https://proceedings.mlr.press/v235/mazeika24a.html.

Microsoft. Fine-tuning safety evaluation, Azure AI Foundry documentation, 2026. Updated July 2026. [CITECHECK: re-check the page].

Ouyang, X., Chen, S., Pearce, M. A. L., Hartvigsen, T., and Schwarz, J. R. ADMIRE-BayesOpt: Accelerated data MIxture RE-weighting for language models with Bayesian optimization. *Transactions on Machine Learning Research*, 2025. URL https://openreview.net/forum?id=0Euvm9zDpu.

Padhi, I., Nagireddy, M., Cornacchia, G., et al. Granite Guardian. arXiv preprint arXiv:2412.07724, 2024. URL https://arxiv.org/abs/2412.07724.

Pan, Y., Shi, T., Zhao, J., and Ma, J. W. Detecting and filtering unsafe training data via data attribution with denoised representation. In *Forty-third International Conference on Machine Learning*, 2026. URL https://openreview.net/forum?id=M9DDNGIM7Z.

Qi, X., Zeng, Y., Xie, T., Chen, P.-Y., Jia, R., Mittal, P., and Henderson, P. Fine-tuning aligned language models compromises safety, even when users do not intend to! In *The Twelfth International Conference on Learning Representations*, 2024. URL https://openreview.net/forum?id=hTEGyKf0dZ.

Qi, X., Panda, A., Lyu, K., Ma, X., Roy, S., Beirami, A., Mittal, P., and Henderson, P. Safety alignment should be made more than just a few tokens deep. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=6Mxhg9PtDE.

Qwen, Yang, A., Yang, B., et al. Qwen2.5 technical report. arXiv preprint arXiv:2412.15115, 2024. URL https://arxiv.org/abs/2412.15115.

Röttger, P., Kirk, H., Vidgen, B., Attanasio, G., Bianchi, F., and Hovy, D. XSTest: A test suite for identifying exaggerated safety behaviours in large language models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 5377–5400. Association for Computational Linguistics, 2024. URL https://aclanthology.org/2024.naacl-long.301/.

Seth, P., Sadhu, S., Kaushal, A., and Sankarapu, V. K. SafeTune: A unified faithful library for auditing and repairing safety drift in fine-tuned LLMs. arXiv preprint arXiv:2609.22153, 2026. URL https://arxiv.org/abs/2609.22153.

Souly, A., Lu, Q., Bowen, D., Trinh, T., Hsieh, E., Pandey, S., Abbeel, P., Svegliato, J., Emmons, S., Watkins, O., and Toyer, S. A StrongREJECT for empty jailbreaks. In *The Thirty-eight Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2024. URL https://openreview.net/forum?id=KZLE5BaaOH.

Springer, M., Lee, C. P., Metevier, B., Castleman, J., Turbal, B., Jung, H., Shen, Z., and Korolova, A. The geometry of alignment collapse: When fine-tuning breaks safety. arXiv preprint arXiv:2602.15799, 2026. URL https://arxiv.org/abs/2602.15799.

Team OLMo, Walsh, P., Soldaini, L., Groeneveld, D., et al. 2 OLMo 2 furious. arXiv preprint arXiv:2501.00656, 2024. URL https://arxiv.org/abs/2501.00656.

Vidgen, B., Scherrer, N., Kirk, H. R., Qian, R., Kannappan, A., Hale, S. A., and Röttger, P. SimpleSafetyTests: a test suite for identifying critical safety risks in large language models. arXiv preprint arXiv:2311.08370, 2023. URL https://arxiv.org/abs/2311.08370.

Wang, M., Xu, Z., Fang, J., Yao, Y., Deng, S., Chen, H., and Zhang, N. From data to behavior: Predicting unintended model behaviors before training. arXiv preprint arXiv:2602.04735, 2026a. URL https://arxiv.org/abs/2602.04735.

Wang, X., Zhang, Y., Liu, Y., Yang, X., Wang, Z., Feng, S., and Wang, D. From parameter dynamics to risk scoring: Quantifying sample-level safety degradation in LLM fine-tuning. In *Forty-third International Conference on Machine Learning*, 2026b. URL https://openreview.net/forum?id=NKWWXlS3xt.

Xia, M., Malladi, S., Gururangan, S., Arora, S., and Chen, D. LESS: Selecting influential data for targeted instruction tuning. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 54104–54132. PMLR, 2024. URL https://proceedings.mlr.press/v235/xia24c.html.

Xiao, F. and Aranguri, S. Probe-based data attribution: Discovering and mitigating undesirable behaviors in LLM post-training. arXiv preprint arXiv:2602.11079, 2026. URL https://arxiv.org/abs/2602.11079.

Yin, J. and Rush, A. M. Compute-constrained data selection. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=4es2oO9tw1.

Zhang, J., Zhou, Q., Deng, X., and Jiang, W. DataRx: Missingness-aware sampling for safer large language model task-specific fine-tuning. arXiv preprint arXiv:2608.04322, 2026a. URL https://arxiv.org/abs/2608.04322.

Zhang, J., Zhou, Q., Deng, X., Jiang, W., Pan, J., and Zhu, J. DataShield: Safety-degrading data filtering for LLM benign instruction fine-tuning. arXiv preprint arXiv:2606.00160, 2026b. URL https://arxiv.org/abs/2606.00160.

Zhao, H., Yuan, C., Huang, F., et al. Qwen3Guard technical report. arXiv preprint arXiv:2510.14276, 2025. URL https://arxiv.org/abs/2510.14276.

A. Editor's notes (not part of the paper; delete before submission)

What this is. Revision 5 of the draft, written on 9 October 2026 from 13:17 UTC (clock checked). It applies the user's decisions of 9 Oct, each checked against FOCUS.md, which wins on any conflict (see "Conflicts", items 19–27): recipe B is the main recipe again and recipe G is the "recipe only" fix and a reference (FOCUS 09:29 UTC); a per-family learning-rate rule replaces the common-recipe ladder (09:29 UTC); five model-level references are mandatory, with Safe LoRA added (10:01 UTC) and the Qi arm on recipe B as B+Qi (13:16 UTC); and the dataset plan is fixed (09:40 UTC). Revision 4 (8 Oct, from 21:40 UTC) withdrew the loss-weighting account of recipe B against recipe G (the token-mean-only run) and added the ARC rule. Revision 3 (from 19:35 UTC) made recipe G the main recipe, which revision 5 reverses, and added the *worse* label. Revision 2 (from 14:00 UTC) follows the paper specification (`research/review/2026-10-08-paper-specification.md`, commit f2a9979) as amended by the user's decision relayed by the lead on 8 Oct: the adopted FOCUS.md amendment (commit 4b37ec2, "ADOPTED 2026-10-08T13:53 UTC") and the study plan `research/review/2026-10-08-sharpened-direction-DRAFT.md` (revision 2). Where the two disagree, the adopted amendment wins. **Nothing in the paper body is a result.** Gates B and C, stage 1, the positive control and the Llama stage have not run; OLMo's Gate A of record (recipe B) stays in this note until §5 is written. Development evidence appears only in this note. Menu sizes, the filter list and costs are marked [TO BE FINALISED]; the Gate B pre-registration (revision 5), the stage-1 pre-registration, Safe LoRA's settings and Alpaca's preconditions are marked [NOT YET FIXED]. The ICML LaTeX version (`icml/main.tex`) carries the same text with a fetched `references.bib`; the author-year citations below use the venue years of those fetched entries.

Title. The spec's working title ends "…, and When Diagnosis Pays". The user's decision makes the measured repair comparison the headline and the decision-value map secondary, so the title now names only the comparison. Alternatives:

1. *Repairing Safety After Benign Fine-Tuning: A Compute-Matched Comparison of Data Repairs, and When Diagnosis Pays* (the spec's title; use it only if the planted sweep is funded and gives the map measured support).
2. *Diagnose or Retrain? A Compute-Matched Comparison of Data Repairs for Safety After Benign Fine-Tuning* (the previous title as a lead-in, as the spec allows).
3. *Choosing a Data Fix When Benign Fine-Tuning Erodes Safety: Diagnosis, Trial Runs and Side Effects at Matched Compute*. This is the DIFFUSE title: the plan (§11) and FOCUS ("agreed fallback") switch to it if Gate B reads DIFFUSE. Its number is kept from the previous draft because the plan refers to it as "alternative title 3".
4. *Remove, Add or Leave the Model Alone? Seeded, Compute-Matched Repairs for Safety Eroded by Benign Fine-Tuning* (stresses the model-level references).

A.1. What changed in revision 5 (9 Oct, from 13:17 UTC)

- **Recipe B is the main recipe again (FOCUS 09:29 UTC, decision 1).** Gate A, Gate B and stage 1 use recipe B (LoRA r64, constant LR 1e-4, no warmup, 128 rows per step, per-row loss, one epoch); OLMo's recipe-B Gate A (PASS, 8 Oct 05:40 UTC) is the Gate A of record. Revisions 3–4's "recipe G main" is undone in the abstract, §1, §4.1, Table 1, §4.3, §4.6 and §6. The abstract and §1 no longer call the main recipe standard practice ("a fixed LoRA recipe that erodes it partly"; "recipe B: a constant learning rate, one epoch"); §6 calls recipe B "plain rather than standard" and says why it was chosen (from probe and Gate A readings, before any repair outcome).
- **Recipe G is the "recipe only" fix and a reference (FOCUS 09:29 and 10:01 UTC).** The body gives its design facts (a 4-step warmup and cosine decay, which halve the mean learning rate, and a token-mean loss) and calls it "the recipe-only fix". Revision 4's token-mean-only finding (the schedule, not the loss weighting, carries G's reduction on OLMo; one seed) stays development evidence; it is why these notes call recipe G a dose-lowering fix. The body does not say that recipe G lowers erosion in every family (conflict 25). "Stronger dose", "stronger-dose comparison" and "lower-dose base" are gone.
- **The per-family learning-rate rule replaces the common-recipe ladder (FOCUS 09:29 UTC, decision 2).** §4.1's "One common recipe" paragraph and [PLACEHOLDER: the common rung] are replaced by "A matched erosion dose": recipe B with only the learning rate changed, the highest rung of (1e-4, 5e-5, 2e-5) whose probe reads partial (20–55 pp) under the frozen probe bars, chosen from probes only, never from a repair result; cross-family claims "at a matched erosion dose". The development ladder (FOCUS 01:43 UTC) is cited, without numbers, as the evidence that one recipe cannot equalize the erosion, and the rule is disclosed as fixed after that ladder was read (§4.1, §6). New marker: [PLACEHOLDER: each family's rung]. Qwen is outside the repair study (§4.1); its recipe-B runs enter only the same-recipe family contrast (§4.10). The contributions, §4.6 (Gate A), §4.8 and §6 say "matched dose".
- **Five mandatory references (FOCUS 10:01 and 13:16 UTC).** The safety prompt; adapter scaling ($\times$ 0.75, $\times$ 0.5); recipe G (3 seeds); B+Qi (Qi et al.'s constrained SFT on recipe B, 6 seeds; "G+Qi" is renamed everywhere); Safe LoRA (Hsu et al., 2024; 3 seeds; settings [NOT YET FIXED]). Changed: abstract, §1, §2 (one clause), §3.2 (now points to §4.4), Table 1 (two rows added), §4.4 (rewritten). New BibTeX `hsu2024safelora`, from the OpenReview venue note HcifdQZFZV (NeurIPS 2024 poster), with arXiv:2405.16833 checked by the arXiv API (CITE-CHECK 26).
- **Dataset plan (FOCUS 09:40 UTC).** §4.1 "Data": Dolly for Gate B and stage 1; PKU-SafeRLHF for the add repairs' safety rows; Alpaca as the second benign dataset after Gate B, [NOT YET FIXED: its units and learning rate]. §6 adds "one benign dataset until Alpaca runs".

- **The two-layer mechanism account (FOCUS 8 Oct 15:17 and 16:50 UTC; the mechanism note's addenda; conflict 19).** §3.6 adds one sentence: the opening-switch term has two layers, the refusal openings lost ($p_0 - p_1$) and how often a direct answer was already unsafe before fine-tuning (u_{N0}). §4.10 item 1 reports both layers across families and recipes; §5's placeholder lists them. Numbers only in the development evidence.
- **§4.10 item 2** is now "Recipe and family contrasts": recipe G against recipe B as a recipe-only fix, the token-mean-only run, the same-recipe family contrast, every learning-rate rung and the rank-16 check under recipe G.
- **HarmBench on the pod and the harm wording (FOCUS 8 Oct 23:43 UTC).** §4.2's last clause is now per recipe ("If the classifier's erosion under a recipe is below 20 pp, no sentence calls that recipe's erosion harm"), so recipe G's erosion is "refusal erosion". Numbers only in the development evidence.
- **ARC rule.** Kept with its disclosure; its scope ("at the common recipe", FOCUS 20:50 UTC) is marked [TO BE FINALISED: its scope under the per-family rule] (decision 16). The sentence "An interval wholly at or below 0.03 passes . . . " is cut (implied by "resolves a straddling interval").
- **Gate B (§4.6).** Its pre-registration, revision 5, is marked [NOT YET FIXED]; "with the warmup and cosine decay spanning its own epoch" is removed (recipe B has no schedule); the exploratory E-scaled floor is removed, because at recipe B it equals the 8 pp floor (decision 21).
- **Order of work (§4.9).** Gate B runs to its verdict; funding for the later stages is then sought with that result; stage 1 starts only once funded, never because of an outcome. The stage-1 pre-registration is marked [NOT YET FIXED].
- **Secondary arm.** "The length rule under recipe B's per-row loss" reverts to revision 2's "the length rule under a token-mean loss" (decision 11).
- **Length.** The body did not grow: the main body measures 9.40 pages against 9.46 for revision 4 by the same measure (pages before the page with the References heading, plus the share of that page's two columns above it, in the ICML 2026 build). To make room, §4.1's common-recipe paragraph, §4.6's E-scaled floor, §4.10 item 2's dose sentence, §4.2's "as before" sentence and §6's repeat of Gate B's power were cut, §4.3's "Re-implementations" paragraph was merged into §6 "Baselines and doses", and Table 1's B+Qi and Safe LoRA rules were kept to one line each.
- **Editor's notes.** Decisions 14 and 15 resolved; decisions 7, 11 and 13 updated; decisions 16–22 added. Risks 2, 3, 17–19, 21–23 updated; risks 24–27 added. Conflicts 9, 10, 12 and 18 marked superseded; conflicts 19–27 added. CITE-CHECK 25 renamed (B+Qi); 26 added. Placeholders, fill-ins, development evidence and status brought up to FOCUS 13:16 UTC.

A.2. What changed in revision 4 (8 Oct, from 21:40 UTC)

History: revision 5 reverses this revision's recipe roles and replaces its common-recipe rule.

- **The loss-weighting account is withdrawn for B → G (FOCUS 21:36 UTC, the token-mean-only run).** The body no longer says or implies that recipe G's token-mean loss, or a first-token loss mechanism, explains why recipe G erodes less than recipe B.
  - §1 and §4.1 describe recipe B as the stronger dose: a constant learning rate without warmup, about twice recipe G's mean learning rate over the epoch (FOCUS: G's warmup and cosine decay "also halves the mean learning rate"), plus a per-row loss. §4.1 drops "which over-weights the first tokens of short answers" and "a mechanism test", and names the token-mean-only run as the run that separates the two changes.
  - §4.10, item 1 is now "refusal openings against harmful content": the decomposition is descriptive, and the first-token margin is the one causal test. The recipe contrast is no longer a causal test. Item 2 replaces "four amplifiers" with the two differences (schedule and loss weighting), adds a placeholder for the token-mean-only run's result against its prediction, and adds that families can respond differently to the same change of dose.
  - §3.6 says the decomposition is descriptive; "should restore less" becomes "may restore less".
  - §5's placeholder list adds the token-mean-only run against its pre-registered prediction, and "the mechanism decomposition and its causal tests" becomes "the opening decomposition and the first-token margin". §6 "Recipe and setting" describes recipe B as a stronger dose and adds the family × dose point and the ARC loosening.
- **G+Qi kept (FOCUS 17:36 UTC, item 4).** §1, §2, §4.4 and Table 1 justify it by Qi et al.'s own objective, which constrains updates on the initial tokens, not by the loss-weighting account; none of that text changed. §4.10 now cites Qi et al. (2025) only for the shallow-alignment observation beside the first-token margin (abstract re-read by the arXiv API between 21:40 and 21:44 UTC; CITE-CHECK 25).
- **ARC rule (FOCUS 20:50 UTC).** §4.2's "Task loss" item becomes "Task loss and skill": the rule in FOCUS's wording, called a loosening, with the disclosure that the first OLMo recipe-G probe was read under the old rule before the rule was proposed. The body gives no numbers for that probe. Revision 3's "ARC . . . sets no limit" was wrong for the probe and Gate A bars, which check ARC; §4.2 now states the bar. §4.1 and §4.6 point to §4.2 for skill; §6 notes the loosening.

- **Common-recipe rule (FOCUS 17:36, 21:04 and 21:37 UTC; the frozen amendment of 7 Oct, 21:35 UTC).** §4.1 adds the too-close rule: one more seed, the two-seed mean classified with the same band, no further seed in that classification, and a choice made after the outcome reported as such. [PLACEHOLDER: the common rung] stays, because Llama's two-seed reading under recipe G is pending and FOCUS records no step-down.

- **Development numbers** (the token-mean-only run, OLMo recipe G at three seeds, Llama under recipe G, the rank-16 check) are only in the development evidence below, labelled. The abstract and contributions did not change.

- **Editor's notes.** Risks 2, 3 and 18 updated; risks 20–23 added; "Where the framing is weak" item 5 updated; decisions 7 and 11 updated, decisions 13–15 added; placeholders updated; CITE-CHECK 25 extended; conflicts 14–18 added; development evidence and status brought up to FOCUS 21:37 UTC.

A.3. What changed in revision 3 (8 Oct, from 19:35 UTC)

History: revision 5 reverses this revision's recipe roles (recipe B is main again) and renames G+Qi to B+Qi.

- **Recipe roles (FOCUS 19:24 UTC; the common-recipe rule, 17:36 UTC).** Recipe G (token-mean loss, 4-step warmup with cosine decay, LR 1e-4, LoRA r64, 128 rows per step, one epoch) is the main recipe for Gate A, Gate B, stage 1 and both families. Recipe B is the stronger-dose comparison; any learning-rate rung that runs is a dose point. Changed: abstract, §1, §3.2, §3.6, §4.1, Table 1, §4.3, §4.4, §4.6, §4.8, §4.10, §5, §6.

- **Llama (FOCUS 16:50 and 17:36 UTC).** "Llama-3.2-1B replicates the comparison" is gone from the abstract, §1, the contributions and §4.8. Llama-3.2-1B is a planned second family under one common recipe, chosen by the pre-registered ladder rule (§4.1). The claim rule "replicates across two families" stays in §4.8 as a rule for when both families have run.

- **Worse label (FOCUS 18:31 UTC).** The labels are better / equivalent / worse / inconclusive (abstract, §1, contributions, §3.2, §4.3). Old decision 2 is resolved.

- **G+Qi reference (FOCUS 17:36 UTC, item 4).** Recipe G plus Qi et al.'s (2025) token-wise constrained SFT, six seeds, a model-level reference with no novelty claim. Recipe G is now the base, so "a standard recipe" leaves the references (FOCUS 17:36 UTC, item 4); the three references are the safety prompt, adapter scaling and G+Qi (abstract, §1, §2, §3.2, Table 1, §4.4). Qi et al. (2025) was already in references.bib (`qi2025shallow`, fetched from OpenReview with arXiv:2406.05946 checked; title and authors re-checked with the arXiv API on 8 Oct between 19:32 and 19:35 UTC; OpenReview now answers with a bot challenge, so the venue was not re-fetched); no new BibTeX.

- **One secondary arm.** "The length rule under a token-mean loss" contrasts nothing once recipe G (token-mean) is the main recipe. §4.3 now names its mirror, the length rule under recipe B's per-row loss. The mirror is the drafter's inference (decision 11).

- **Editor's notes.** Development evidence and status brought up to FOCUS 19:24 UTC; risks 2 and 3 updated, risks 18–19 added; conflicts 9–13 added.

A.4. What changed in revision 2 (from the draft of 8 Oct, about 12:00 UTC)

- **Headline.** The measured repair comparison leads (six repairs × six seeds against random removal and no repair, labels better / equivalent within 5 pp / inconclusive, with over-refusal, task loss and measured cost). The "when diagnosis pays" map is secondary and model-based; its predictions are written down before Gate B's category runs and tested on Llama.

- **Menu.** Stage 1's primary family P1–P6 (A1, L1, H1 removals; S1, S3, S4 additions); secondary arms and scored-only scorers; rewrite and self-distill only in the DIFFUSE branch. The previous draft's keep/remove/add/rewrite menu by source, the default mix a_{mix} and the "beat a_{mix}" rule are gone.

- **New parts.** Three mandatory references with no data change (safety prompt, adapter scaling, recipe G; revision 3 replaces recipe G with G+Qi); the planted-cause positive control; the HarmBench claim rule; two cost conventions; the Llama rule (LR or rank only; revision 3: the common-recipe ladder, LR only, for both families); confirmation on a sealed held-out subsample; the opening decomposition (side study).

- **Resolved old decisions.** O is measured against the untuned model (old decision 1); τ_O is a grid of 2, 5 and 10 pp and τ_L is 2 SE of a seed-to-seed difference in held-out loss on all 200 rows (old decision 2); the main scale is 1B with a ladder (old decision 3); a judge rule decides "harm" against "refusal" wording (old decision 4); model-level references are mandatory (old decision 6); side-effect prediction is secondary, on the add arms only (old co-primary).

- **§4 and §6 are written out. §2** adds the published filters, the format paper, DualSelect, Bianchi et al., PTST and MNEME, the DataRx correction and Guo et al.'s Table 2.

A.5. Placeholders

Where	Marker	Filled by
Abstract	[PLACEHOLDER: if funded]; [FINDING 1]–[FINDING 3]	the funding decision for the planted sweep; stage 1, the references, Gate B and the sweep
§1	[PLACEHOLDER: if funded]; [TO BE FINALISED] menu sizes 8/32/128	as above; the plan's review
§1, contributions	[FINDING 1]–[FINDING 4]; [ARTIFACT]	stage 1 and Llama; the sweep; Gates B/C and the mechanism study; the decision-value predictions; the anonymized outcome tables
§1	[FIGURE 1]	the repair-comparison figure (restoration with 90% intervals and the $\pm$5 pp band against over-refusal and cost, with the adapter-scaling curve)
§3.2	[TO BE FINALISED] the label when a repair is better on U but worse on a side effect	the stage-1 pre-registration (decision 12)
§3	the constants δ (5 pp, given), δ_B (8 pp, given), τ_O grid (given), τ_L (rule given, value after Gate A noise), B, the selector rules per (M, B), k_d, ϕ_j and the V level	the stage-1 pre-registration and the selector spec (committed before Gate B's category runs)
§4.1	[PLACEHOLDER: sizes actually run]; [NOT YET FIXED: Alpaca's units and learning rate]; [PLACEHOLDER: Llama-3.2-1B's rung]	funding (scale ladder); the Alpaca design note and the user's extension of the learning-rate rule to datasets (decision 20); OLMo: rung 1, recipe B at 1e-4 (its probe and Gate A of record read partial); Llama: its recipe-B probes under the rule (decision 18)
§4.2	[TO BE FINALISED: the ARC rule's scope under the per-family rule]	the user (decision 16)
§4.3, Table 1	[TO BE FINALISED] secondary arms and the filter list; [PLACEHOLDER: final MDE]; [NOT YET FIXED: Safe LoRA's settings]	the plan's review; Gate B's pooled noise; the stage-1 pre-registration (decision 19)
§4.4	[NOT YET FIXED] Safe LoRA's base model, layer rule and threshold	the stage-1 pre-registration, after the lead reads the paper in full (decision 19)
§4.5	[TO BE FINALISED] menu sizes; [PLACEHOLDER: runs only if funded]	the plan's review; funding
§4.6	[NOT YET FIXED] Gate B's pre-registration, revision 5	the user's acceptance of DRAFT revision 5 (decision 7)
§4.7	[TO BE FINALISED] costs	measured ledgers; frozen cost conventions
§4.9	[NOT YET FIXED] the stage-1 pre-registration	committed before Gate B's category runs
§4.10	[PLACEHOLDER: report against the recorded bound]; [PLACEHOLDER: its result against its pre-registered prediction, §5]	the mechanism prediction recorded before Gate B; the token-mean-only run (read; development evidence), entered in §5 when results are written
§5, §7	placeholders	results
References	Microsoft page [CITE-CHECK]; Alpaca has no citation (§4.1, §4.8)	the author

[TO BE FINALISED] items (the plan is under review): (1) the planted menus' sizes (8, 32, 128 cells) and cell rules; (2) the published-filter list: L1 is LARF-style in the primary family; DataShield-style is secondary; Self-Inf-N-style is scored only; SQSD is cited only; bi-anchoring enters as A1's score; (3) all costs: the FRE values, both cost conventions, the cost of each scorer and of A1, the lean-trial cost c_0, and the dollar plan (stage 1 about USD 56–98; planted sweep about USD 17–35; core path about USD 215–405; all estimates from the plan, not paper text); (4) the secondary arms and their order.

Fill-ins for each outcome (templates, not claims):

- **A repair is better.** F1: "[repair] restores __ pp more than random removal (90% CI a–b), with over-refusal __ pp above the untuned model and task loss within τ_L, at __ FRE; it [does / does not] lie below the adapter-scaling curve". If it does not, the sentence must say that a model-level reference does as well (FOCUS "do not claim": data repair beats model-level repair).
- **A repair is worse.** F1: "[repair] is worse than [comparator]: it [raises the unsafe rate / raises over-refusal / raises task loss] by __ (Holm-adjusted p = __)". This label replaces "inconclusive" for any significantly harmful contrast (FOCUS 18:31 UTC); P4 (S1 against no repair) is the contrast where prior work predicts a harmful result (DataRx's random mix).
- **Everything equivalent.** F1: "every removal repair is equivalent to random removal within $\pm$5 pp; [addition result]". This is a publishable answer; the abstract then says it.

- **A free reference wins.** F2: "a safety prompt / a down-scaled adapter / Safe LoRA / recipe G / B+Qi restores as much as any data repair at lower over-refusal" (plan §13: a valid, reportable answer).
- **Planted cause found / not found** (only if funded). F3: "[diagnostic] ranks the planted cell first at M = __ ; runs alone find it at __ FRE".
- **DIFFUSE.** F3: "no Dolly category restores more than 8 pp beyond random removal (80% detection at about 10–12 pp)". Switch to alternative title 3; the planted sweep becomes the only test of whether diagnosis can pay.
- **Second family.** If Llama's Gate A passes at its matched dose and stage 2 runs: "[contrast] replicates on Llama-3.2-1B at a matched erosion dose (same sign; family difference below the interaction's MDE of __ pp)". If no learning-rate rung of recipe B reads partial for Llama, or stage 2 is not funded, the paper is single-family: drop the second family from the abstract and contributions, and say so in §6.
- **Harm wording.** Use "restored refusal" unless the HarmBench rule (§4.2) passes for that contrast; never "harm" if E_{HB} ¡ 20 pp. Recipe B's OLMo erosion clears 20 pp on every seed and recipe G's does not (development evidence), so recipe G's erosion is "refusal erosion under Qwen3Guard", never harm.

A.6. [CITE-CHECK] items

#	Item	Where	Why
1	TOSS (2603.01185)	§2	Title and venue now fetched ("Token-level Data Selection for Safe LLM Fine-tuning", ICLR 2026). The claim "token-level selection" rests on the title only.
2	Unforgotten Safety (2512.10150)	§2	Fetched ("... with Continual Learning"); claim from the title only.
3	Yin & Rush (2410.16208, ICLR 2025)	§1, §2	Wording ("almost never compute-optimal"; 5x and 10x model-size ratios) from the arXiv abstract, read this session. The spec's "predicts retraining wins at equal size" is our inference from the 10x ratio; check the paper before using it.
4	TuneAhead (2606.17660)	§2	Fetched; claim from the title only.
5	Data2Behavior (2602.04735)	§2	Fetched; claim from the title only.
6	SafeTune: model-level repairs, single deterministic runs	§2	Recipe and runs PDF-checked (benign-erosion recipes note); "compares repairs of safety drift" from the title and a fetch summary.
7	DataRx mechanism ("gap between the model's own response and a safe reference")	§2	From the arXiv abstract (this session). The numbers 59.23% and 13.70% are in the abstract; 20.24% is Table II (recipes note, PDF read); "averaged over seven tasks" is the abstract's wording. The plan reads (by HTML fetch, not PDF) that DataRx's random arm draws refusal-only rows; that is why S1 is the DataRx test.
8	DataShield and GradShield report scoring cost without charging it against runs; DataShield's mechanism	§2	Cost claim: fetch summaries only. Mechanism (compliance direction): DataShield's abstract (this session). DataShield's cost table (Table IV) is cited by the plan from the reading list (LARF 0.17 h, DataShield 0.50 h, bi-anchoring 3.5 h per 10,000 rows); not PDF-checked.
9	SOURCE "reported to cost several training runs"	§2	Reported through the Bergson paper; not checked.
10	LESS as "gradient similarity on LoRA gradients"	§2	Reading-list summary.
11	Persona vectors	§2	Blog post via the problem-validation review; confirm in 2507.21509.
12	DataDecide; Ivison et al.	§2	Positioning review, from an extraction-tool read.
13	Xiao & Aranguri	§2	Abstract only ("over 10 times cheaper").
14	Microsoft fine-tuning safety documentation	§1, refs	Read on 6 Oct; re-check before submission; no URL in the bib.
15	Venues and author order	refs	ADMIRE-BayesOpt: TMLR lists Ouyang first, arXiv lists Chen first (bib follows TMLR). DATE-LM year set to 2025 from the venue string.
16	LARF (2507.18631, EMNLP 2025)	§2	Mechanism from the abstract (this session).
17	SQSD (2605.04572, ICML 2026)	§2	Mechanism and acronym from the abstract (this session). "Its scoring needs several FRE" is the plan's estimate (6–14 FRE), not the paper's.
18	Self-Inf-N (2505.06843, ICML 2025)	§2, §4.3	"100 outliers break safety; most mitigations fail" from the abstract (this session). Using its score as a removal filter is our adaptation.
19	"Bi-anchoring" as He et al.'s method name	§2, §4.3	The reading list's name; check He et al.'s own term.
20	Form over content (2609.19589)	§1, §2	Abstract only (here and in the repo's arXiv search file).

#	Item	Where	Why
21	DualSelect (2606.09866)	§1, §2	Abstract (this session): joint task and reference selection, 1B–8B models. "Compares random, SafeInstr, SEAL, LARF, TOSS and SQSD with three seeds and a false-refusal rate" comes from the plan's novelty review (HTML §4), not re-read.
22	Bianchi et al. (ICLR 2024): 3% safety data helps, with exaggerated safety	§2	Abstract (this session).
23	MNEME (EMNLP 2025) predicts side effects	§2	Abstract (this session).
24	Instrument citations (OLMo 2, Llama 3, Qwen2.5, Qwen3Guard, Granite Guardian, HarmBench, StrongREJECT, SimpleSafetyTests, XSTest, OR-Bench, PKU-SafeRLHF, LoRA, ARC, Dolly)	§4	Fetched. Check that each cited report covers the version used (Granite Guardian 3.3 8B; Qwen3Guard-Gen-4B; Llama-3.2 under the Llama 3 report; OLMo-2-0425-1B under the OLMo 2 report). Dolly's entry is the dataset card's own BibTeX (a blog post). Alpaca (§4.1, §4.8) has no citation yet: the brief allows new BibTeX only from the arXiv API, OpenReview or PMLR, and Alpaca was not searched there this revision (it is a GitHub release).
25	Qi et al. (2025; 2406.05946) as the B+Qi reference (renamed from G+Qi in revision 5; "token-wise constrained SFT", "constrains updates on the first answer tokens")	§1, §2, §3.6, §4.4, §4.10	Mechanism from the arXiv abstract, read by API this session ("constraining updates on initial tokens"). "Token-wise constrained SFT" is the lead's and FOCUS's name; check the paper's own term before submission. FAQ (FOCUS 15:59 UTC) reports that their Table 9 shows warmup alone does most of the work on benign data; not checked here. Revision 4: §4.10 also cites them for "safety alignment often adapts mainly the first few output tokens", a paraphrase of the abstract ("the alignment adapts a model's generative distribution primarily over only its very first few output tokens"), re-read by the arXiv API on 8 Oct between 21:40 and 21:44 UTC.
26	Safe LoRA (Hsu et al., 2024; arXiv:2405.16833; NeurIPS 2024) as a reference ("projects the LoRA weights of selected layers onto a safety-aligned subspace without training or data")	§1, §2, §4.4, Table 1	BibTeX from the OpenReview venue note HcifdQZFZV (NeurIPS 2024 poster), fetched by the OpenReview API on 9 Oct at about 13:20 UTC; title and the six authors checked against the arXiv API record (v2, "camera-ready version accepted for NeurIPS 2024"). The body's wording is from the arXiv abstract, read by API this session ("projection of LoRA weights from selected layers to the safety-aligned subspace"; "training-free and data-free"; it "only requires the knowledge of the weights from the base and aligned LLMs"). FOCUS 10:01 UTC's fuller description (the weight difference between the aligned and the base model; the layers where the update departs from it) was not checked against the paper; the lead reads it in full before the settings are fixed (decision 19).

PDF- or full-text-checked in the repo, and used as such: Qi et al. (2024), He et al., Luo et al. (2026b), Pan et al., DATE-LM (verification log, 6 Oct); Kim et al. (2025), SafeTune's recipe, PTST, DataRx's Table II (benign-erosion recipes note); Guo et al. including Table 2 (7.27, 4.50, 9.09, 44.80 confirmed), Springer et al. (safety-landscape note); Llama 3's safety mix (problem-validation review); Qi et al. (2025) on first-token alignment (erosion-amplifiers note). Guo's "two seeds" is the Table 2 caption; the Guo-gap note finds no seed-69 checkpoint for the 5,000-row cells.

A.7. Decisions the author must make

1. **Title** (above).
2. **RESOLVED (FOCUS 18:31 UTC, the user): the "worse" label.** A repair whose contrast against its comparator is significantly harmful at the 5% Holm level (on unsafe share, over-refusal or task loss, each in its harmful direction) is labelled *worse*, so a harmful repair is never reported as inconclusive. Labels are better, equivalent, worse and inconclusive; §3.2 and §4.3 carry them, and P4's harmful direction (DataRx) now has a label. One detail is left for the stage-1 pre-registration (decision 12).
3. **The planted sweep's funding.** If it is not funded, "menu size" leaves the headline: delete contribution 2, the planted part of the abstract and Figure 1's menu-size panel, and §3.4 shrinks to one sentence.
4. **Preliminary-results section.** The spec does not ask for one, so the body has none. Gate A's numbers and their disclosures stay in this note until §5 is written. Decide whether a section labelled "Preliminary results (development evidence; not confirmatory)" should carry Gate A.
5. **Filter list and secondary arms** ([TO BE FINALISED]).
6. **Venue and length.** The adopted amendment targets TMLR or NeurIPS 2027; ICML 2027 only if, by about 1 December, the planted sweep, both 1B families and one 7–8B published-setting replication are funded. The LaTeX build uses the ICML 2026 style; see

the page count in the build log. NeurIPS allows 9 pages; TMLR has no strict limit. ICML needs an Impact Statement (missing). **Revision 5:** the main body measures 9.40 pages (the measure in "What changed in revision 5"), over ICML's 8 pages and NeurIPS's 9.

7. **The Gate B and Gate C rules** in §4.6 summarize the DRAFT pre-registration (revision 4b, written for recipe B), which awaits the user's acceptance. FOCUS 17:36 UTC moves Gate B to the common recipe (REC and recipe S rebuilt from it; the noise gate and predictions P3–P4 re-based on that recipe's Gate A mean E; the 8 pp floor unchanged; each removal arm's warmup-cosine schedule spans its own epoch); DRAFT revision 5 is to carry this and was not yet in the repository at this revision. §4.6 states only the recipe change. Update §4.6 when revision 5 is accepted. **Revision 4:** the repository's DRAFT is still revision 4b. FOCUS 20:17 UTC adds an exploratory, report-only E-scaled floor for Gate B, $8 \times E_{\text{common}}$ / 45.1 pp (about 5.2 pp at E $\approx$ 29), labelled exploratory wherever it appears and never changing a verdict; the 8 pp floor stays primary. §4.6 does not mention it yet; add one sentence when revision 5 is accepted. **Revision 5:** FOCUS 9 Oct 09:29 UTC keeps Gate B on recipe B, so the move to the common recipe above is void; DRAFT revision 5 "stays on recipe B and adds the [D14]–[D19] rules", and the brief's [D20] fixes the kit's price constants (no hourly limit; guard planning rate USD 1.85/h; spend rate USD 1.79/h; reserve USD 6), with the money thresholds recomputed in revision 5 (FOCUS 13:16 UTC). The repository's DRAFT is still revision 4b, so §4.6 marks revision 5 [NOT YET FIXED]. Update §4.6 when revision 5 is accepted.

8. **The opening decomposition's definition** (§3.6) must match the exploratory note's "composition share"; the note does not state its formula in these terms. Fix the order of terms and report both orders.

9. **Gate names.** Keep "Gate A/B/C" or call them pre-registered checks.

10. **Who the user is** (§6): the plan names a fine-tuning service that must ship the customer's adapter and may change only data, or a data-only pipeline. Confirm this wording.

11. **The length-rule secondary arm under the other loss** (§4.3). Revision 2 listed "the length rule under a token-mean loss", a control for recipe B's per-row loss. With recipe G (token-mean) as the main recipe, revision 3 writes the mirror: the length rule under recipe B's per-row loss. Confirm, or drop the arm. **Revision 4:** the token-mean-only run (development evidence) finds that the loss weighting carries none of recipe G's reduction at one seed, so the arm's reason (that per-row weighting makes short rows matter more) has no support at the recipe level. The arm would still test it at the row level; keep it only if three seeds are cheap. **Revision 5:** with recipe B (per-row loss) the main recipe again, §4.3 reverts to revision 2's "the length rule under a token-mean loss". Keep it only if three seeds are cheap, or drop it.

12. **Label order when endpoints disagree.** The 18:31 UTC rule labels a repair *worse* if any of its three contrasts (unsafe share, over-refusal, task loss) is significantly harmful. It does not say which label a repair gets when it is *better* on the unsafe share and *worse* on over-refusal or task loss. §3.2 reports both facts and leaves the order to the stage-1 pre-registration; fix it there.

13. **Where the failed prediction goes.** FOCUS 21:36 UTC: "the paper reports it as a pre-registered prediction that failed". Revision 4 puts its result in §5 (placeholder) and names the run in §4.1 and §4.10. Decide whether §5 gives it a sentence or a row in the dose table, and whether the mechanism side study keeps any recipe contrast. Revision 4 keeps the recipe contrast only as a dose decomposition, not as a mechanism test. **Revision 5:** recipe G is now a reference arm, and the token-mean-only run explains that reference's effect on OLMo; §4.10 item 2 and §5 keep it. Decide whether §5 gives it a sentence or a row beside recipe G.

14. **RESOLVED (FOCUS 9 Oct 09:29 UTC): no common recipe.** Llama's two-seed reading under recipe G stayed TOO CLOSE (+58.25 pp), and the ladder found no rung for both families (01:43 UTC); the user chose recipe B as the main recipe and the per-family learning-rate rule (§4.1). Llama's recipe-G seed 1103 was then read as a report-only third seed (three-seed mean +57.1 pp, TOO CLOSE at the top edge; results record, 09:35 UTC).

15. **RESOLVED (FOCUS 9 Oct 09:29 UTC): recipe G's name.** Recipe G is "the recipe only fix" and a reference, not the main recipe or a lower-dose base; the body calls it "the recipe-only fix", and the abstract no longer calls the main recipe standard practice.

16. **The ARC rule's scope under the per-family rule.** FOCUS 20:50 UTC adopted it "for the probe and Gate A readings at the common recipe (both families)", and no common recipe now exists. OLMo's Gate A of record (recipe B) passed ARC under the frozen bars (ARC rose on every seed; audit), so the rule decided none of OLMo's main readings; it decided the recipe-G Gate A read and the rung-2 readings (reported). Decide whether it applies to Llama's probes and Gate A under the per-family rule; §4.2 is marked [TO BE FINALISED].

17. **The family $\times$ recipe interaction test (Side Project).** Its 95% interval on the logit scale excludes 0, but it was written after 11 of its 12 values were known, has about 5 df and uses one judge (development evidence). Decide whether §5 reports it, labelled exploratory, or only the per-family rule it motivates. The body now gives only the design consequence (§4.1: "families respond to the same change of recipe by different amounts").

18. **Llama's rung under the per-family rule.** Llama's rung-1 probe at recipe B (8 Oct, 16:50 UTC) read SATURATED, and the two report-only seeds agree (three-seed mean +59.3 pp), so the rule's next step would be a recipe-B probe at 5e-5. This is the drafter's reading: FOCUS records no such run or decision. The rung-2 reading under recipe G (+41.2 pp, partial) is not a recipe-B probe and cannot set Llama's rate. [PLACEHOLDER: Llama-3.2-1B's rung] in §4.1 waits for it.

19. **Safe LoRA's settings** (FOCUS 10:01 UTC; to fix in the stage-1 pre-registration before any Safe LoRA run, after the lead reads the paper in full): the unaligned model (OLMo-2-0425-1B base for OLMo-2-0425-1B-Instruct; Llama-3.2-1B base for the second family), the exact projection, the layer-selection rule and its threshold (from the paper's rule or a fixed default, never tuned on stage-1 outcomes), the seeds (the adapter-scaling seeds 1101–1103 unless the pre-registration says otherwise; Table 1 says 3), and the engine work (a post-hoc state transform evaluated by the existing eval-state job, built with tests and reviewed). §4.4 and Table 1 are marked [NOT YET FIXED].

20. **Alpaca's preconditions** (FOCUS 09:40 UTC; to settle before any Alpaca run, not now): the user extends the learning-rate rule from families to datasets before any Alpaca probe; a short design note fixes Alpaca's units (it has no categories, unlike Dolly), the field mapping (instruction, input as context, output as response), the version (tatsu-lab/alpaca; CC BY-NC 4.0) and the download, which the user approves. Alpaca still has no citation (CITE-CHECK 24).

21. **The exploratory E-scaled Gate B floor** (FOCUS 8 Oct 20:17 UTC): $8 \times E_{\text{common}}$ / 45.1 pp equals 8 pp at recipe B (E_{common} is the recipe-B Gate A mean, +45.1 pp), so it repeats the primary floor; revision 5 drops the sentence from §4.6. Confirm, or restore it if DRAFT revision 5 keeps it.

22. **B+Qi's warmup.** FAQ reports that Qi et al.'s Table 9 shows warmup alone does most of the work on benign data (FOCUS 8 Oct 15:59 UTC; not checked here). Recipe B has no warmup, so whether B+Qi takes Qi et al.'s own schedule decides whether B+Qi differs from recipe B only in the objective or also in the schedule, like recipe G. Fix it in the stage-1 pre-registration and report it beside recipe G. **Lead's default (FOCUS 13:45 UTC, 9 Oct):** Qi's constrained loss on recipe B *without* Qi's warmup, so B+Qi isolates the loss while recipe G covers the schedule; the stage-1 pre-registration fixes it [NOT YET FIXED].

A.8. Risks a reviewer will raise (sources in brackets)

1. **"Unsafe" means "did not refuse".** See the development evidence. Answer: the HarmBench claim rule, hard-category U, U by length band, the selectivity score, and claim wording [Gate A audit, 4a; plan §4.2].

2. **Forking paths.** Five settings were screened before OLMo-2-1B fell in the 20–55 pp window, fixed after the Qwen results; seed 1101 both chose the setting and counts in the mean [audit, 4b]. Answer: report every setting tried and the fresh-seed mean (+45.9 pp). **Added in revision 3:** recipe G became the main recipe after its one-seed OLMo result (+29.0 pp, partial) and Llama's saturation under recipe B were known (FOCUS 17:36 and 19:24 UTC). The choice used probe data only, never a Gate B or stage-1 outcome; Gate A is re-run at the common recipe with seeds 1102 and 1103 per family, and seed 1101 (the probe) counts, so report the fresh-seed mean beside the 3-seed mean again. **Added in revision 4:** the ARC rule was adopted (FOCUS 20:50 UTC) after the first OLMo recipe-G probe had been read under the old rule (disclosed in §4.2); OLMo's recipe-G seeds 1102 and 1103 ran before Llama's recipe-G probe was read and count toward Gate A only if recipe G becomes the common recipe (19:58 UTC); Llama's recipe-G seed 1103 was launched before seed 1102 was read and never enters the two-seed classification (21:37 UTC). **Added in revision 5:** the recipe roles changed twice on development data: recipe B (Gate A, 8 Oct morning), recipe G main (8 Oct 17:36 and 19:24 UTC), recipe B main again (9 Oct 09:29 UTC), after the ladder and the HarmBench pod results were known. Every switch used probe, Gate A and judge readings only, never a Gate B or stage-1 outcome, and the Gate A of record is the recipe-B one, read before any of the switches. The per-family rule was fixed after the ladder was read; §4.1 and §6 say so.

3. **Recipe dependence.** On OLMo, recipe G erodes about a third less than recipe B, and the token-mean-only run puts all of that on the schedule (development evidence, one seed). Answer (revision 5): recipe B is the main recipe and the body calls it plain, not standard; recipe G, the standard schedule, is a mandatory reference (the recipe-only fix), so the reader sees what the schedule alone removes; the repair contrasts are within recipe B, and §6 says rankings may differ at a lower dose. On Llama recipe G barely moves the erosion (risk 26), so "G lowers the dose" is a statement about the learning rate, not about erosion in every family. The rank-16 check under recipe G reads "rank matters little" by its pre-registered 10 pp rule, but with a small margin at one seed each, and over-refusal rose at rank 16 (development evidence); report the difference itself, not only the rule's label (FOCUS 21:25 UTC). The rank-64, all-linear adapter is still larger than common practice (rank 8–16).

4. **Nulls everywhere.** A 600-row filter clears the MDE only if the findable effect is about 11–14 pp at row AUC 0.9 [plan §1 item 16, POW Part 3]. Answer: equivalence labels make nulls informative; the planted control; the dose curve.

5. **Additions win on U by construction.** Refusal rows restore refusal, and U tracks refusal [plan §13]. Answer: the HarmBench rule, selectivity, O against the untuned model, and the adapter-scaling curve.

6. **A free reference may beat every data repair** (safety prompt, adapter scaling) [plan §13]. That is a valid answer; say so.

7. **Prior-work overlap** (DualSelect, Yin & Rush, DataShield's cost table). Answer: no "first" claims (§2 "How we differ"); the equivalence bounds, references, positive control and pre-registered predictions are the contribution.

8. **Re-implemented baselines.** The four method repositories have no licence [plan §5.1, FIL]; all scorers are re-implemented ("X-style") with deviation tables; authors are asked for reference outputs; stage 4's published-setting replication checks fidelity.

9. **Diffuse erosion is the expected Gate B outcome** [Gate B draft §8; audit §7]. Under this framing it no longer collapses the paper: stage 1 runs unchanged, and the planted sweep carries the diagnosis question.

10. **The plant may not erode** above the background [plan §13]; the pilot gate decides.

11. **Small scale favours trial runs; cost ratios change with size.** No 1B cost ratio is claimed for 7–8B.

12. **The decision-value map is a model output** at the menus we can afford [plan §0, NOV-2, MET-M1]. Answer: it is secondary, frozen before the data, and scored against measurements and on Llama.

13. **Judge precision on benign prompts.** After fine-tuning Qwen3Guard labels 22–29% of the benign bank Unsafe (untuned 5%) [audit §5].

14. **"Adding a little safety data is known to work." Correction (FAQ second review, 8 Oct; source opened):** DataRx does NOT show that adding works by default. Its 59.23% $\rightarrow$ 13.70% compares a *random* 1% BeaverTails mix with DataRx's *selected* 1% mix on Llama3-8B. The fine-tune with no safety data scored 20.24%, so the random mix made it worse. So *which* safety data is added decides whether adding helps (now P4 and P5). Guo's 10 rows were not checked for over-refusal.

15. **Added data can hide misalignment** that returns on prompts like the training data (Dubiński et al.) [reading list v2, risk 4]. The plan adds an overlap control: the add effect must hold on the eval prompts least similar to the add set.

16. **Motivation.** We found no public evidence that practitioners weigh attribution against retrains [problem-validation review]. The draft does not claim that they do.

17. **Money.** The balance is USD 57.46 (FOCUS 9 Oct 12:18 UTC) against a core path of about USD 215–405 [plan §10, §13]. Gate B's wave-1 need is being recomputed with the [D20] constants (a rough estimate is about USD 60–64, so a top-up of about USD 13–17; the recomputed figure decides; FOCUS 13:16 UTC). Every stage starts on a funding event; only complete seed blocks are analysed.

18. **B+Qi may add little over recipe B, or act like recipe G.** FAQ reports that Qi et al.'s Table 9 shows warmup alone does most of the work on benign data [FOCUS 15:59 UTC; CITE-CHECK 25], and our token-mean-only run puts recipe G's reduction on its schedule (development evidence). If B+Qi brings Qi et al.'s warmup to recipe B, it differs from recipe B in its schedule as well as its objective (decision 22). B+Qi stays as a reference because Qi et al.'s objective constrains updates on the initial tokens, a mechanism different from loss weighting; the paper must not justify it by the refuted loss-weighting account. A B+Qi equal to recipe B or to recipe G is still a valid reference result; no novelty is claimed for it.

19. **One family, or a same-recipe contrast read as a replication.** Under recipe B, Qwen and Llama saturate or sit at the top edge and OLMo erodes partly. That is the same-recipe family contrast (§4.10), reported separately and not as a replication. Under the per-family rule Llama runs at its own rate; if no rung of recipe B reads partial for Llama, or stage 2 is not funded, the paper has one family.

20. **The mechanism account is refuted for B $\rightarrow$ G.** The mechanism note's first-token loss-mass account (under the per-row loss the first 5 response tokens carry 17.7% of the objective, against 5.3% under the token mean; FOCUS 20:01 UTC) predicted that the token-mean loss carries most of recipe G's reduction. The pre-registered token-mean-only run found that it carries none, at one seed [FOCUS 21:36 UTC]. Answer: report it as a pre-registered prediction that failed; the body offers no loss-weighting explanation; the side study is "refusal openings against harmful content", descriptive, with the first-token margin as its only causal test. Under recipe G refusal openings fall less than under recipe B, but that is now read as a dose effect, not as support for the loss account.

21. **Recipe G is a lower dose.** A reviewer may read "the standard schedule erodes less" as "a lower mean learning rate erodes less", which work on optimization choices already suggests (Kim et al., 2025). Answer: say so; recipe G is a recipe-only reference, and the main claims are repair contrasts within recipe B.

22. **Family $\times$ recipe interaction.** Recipe G cuts OLMo's erosion by about a third but barely moves Llama's, and the Side Project's logit-scale test excludes 0 (exploratory; development evidence). No rung of the one-recipe ladder left both families partly eroded (FOCUS 01:43 UTC). Answer (revision 5): the per-family rule matches the erosion dose, and cross-family sentences say "at a matched erosion dose"; the rule was fixed after the ladder was read (disclosed in §4.1 and §6); a family difference is claimed only above the interaction's MDE.

23. **The ARC rule is a loosening, adopted mid-stream.** It can pass a run whose true drop is up to about 6 points, and for the first OLMo recipe-G probe it was not applied blind [FOCUS 20:50 UTC]. Answer: both verdicts are reported; "skill kept" rests mainly on the held-out task loss; §4.2 and §6 disclose it. **Revision 5:** its scope after the per-family rule is open (decision 16).

24. **The main recipe is harsh.** Recipe B has no warmup, a constant learning rate and a per-row loss; a reviewer may call the erosion an artefact of a recipe few people use. Answer: recipe G, the standard schedule, is a mandatory reference, so the paper shows what the schedule alone removes; recipe B was chosen because its erosion clears the 20 pp HarmBench line (so harm wording is allowed) and leaves more room for repairs, while recipe G's does not (development evidence; FOCUS 09:29 UTC); §6 says repair rankings may differ at a lower dose.

25. **Safe LoRA may beat every data repair.** It needs no training and no data. That is a valid answer, like risk 6, and the paper says so. Its settings must be fixed before any stage-1 outcome (decision 19).

26. **"G lowers the dose" does not carry over between families.** Halving recipe G's peak learning rate (rung 2) cut both families by about 28–29%, but moving from recipe B to recipe G, which halves the mean learning rate, cut OLMo by 36% and Llama by only about 4–7% (Llama's B-against-G gap is within its seed spread; development evidence). So the schedule does not act like a lower constant rate in every family, and the token-mean-only run is OLMo-only and one seed. Answer: the body states only the design fact (G halves the mean learning rate) and calls G the recipe-only fix; it makes no claim that G lowers erosion in every family.

27. **A matched dose differs in learning rate.** Under the per-family rule a contrast that "replicates at a matched erosion dose" replicates across families and learning rates together. Answer: §4.8's claim rule and the fill-in say "at a matched erosion dose"; the same-recipe contrast is reported separately.

A.9. Where the framing is weak (the drafter's view)

1. **The cost axis is mostly evaluation.** Generation is about 81% of an arm and training about 8% [plan §1 item 6]. FRE then measures our evaluation size more than training; the two cost conventions help, but a reviewer may still read "compute-matched" as "evaluation-matched". Say plainly what one FRE buys.

2. **The comparison fixes doses.** Rankings across removal (600 rows) and addition (48 rows) hold only at these doses; the dose curve is secondary and 3-seed.

3. **Category units versus row scorers.** Gate B tests categories; stage 1 tests row scorers. A row scorer can find what Gate B cannot, and the reverse; the paper must not read one from the other.

4. **The data-only user and diagnostics pull against each other** (previous draft, point 2): a closed-service customer cannot run attribution; a developer with open weights also has model-level repairs. The references now make this explicit, but the population that benefits from paid diagnosis stays narrow.

5. **Mechanism tests are thin.** The token-mean-only run separated loss weighting from schedule at one seed, on OLMo only, and refuted the loss-weighting account of B $\rightarrow$ G. What remains is a descriptive decomposition with its two layers (openings lost; direct answers already unsafe) and one teacher-forced first-token margin. The two-layer account is observational (no prefill test; positions counted in words), and neither part says why fine-tuning switches refusal openings to direct ones. Both are side studies.

A.10. Development evidence (for the author only; never a finding in the abstract or the contributions)

- **Gate A PASSED on OLMo-2-0425-1B-Instruct under recipe B** (8 Oct, 05:40 UTC; the Gate A of record since FOCUS 9 Oct 09:29 UTC): E = +43.5, +43.7, +48.2 pp over three seeds; mean +45.1 pp (band 41.3–49.0; partial); fresh-seed mean +45.9 pp; answers end; held-out Dolly loss 2.10 $\rightarrow$ 1.54; Granite agrees. Independent audit: sound with caveats.

- **Audit disclosures:** P(Unsafe — no refusal) 0.91–0.94 after fine-tuning; about a third of Unsafe answers carry only "unethical acts" or "politically sensitive"; E under hard categories only is about 28 pp; five settings screened; 2 df for seed variance; the ARC bar is near chance at 1B (untuned 0.392).

- **Qwen2.5-1.5B and 7B-Instruct saturate** under recipe B (+75.2 and +82.2 pp); rank 16 still gives +67.8 pp. Guo-gap analysis: judge and prompts explain at most about 13–30 pp of the gap to Guo's Qwen LoRA cell.

- **Mechanism (exploratory note):** refusal openings fall by about 50 pp (OLMo) and 85 pp (Qwen); the opening switch carries about 0.42–0.64 of OLMo's Unsafe increase; direct answers also became more harmful; the erosion is broad across prompts and stable across seeds. Recorded prediction: a first-token path alone gives at most about 3 pp per category, so a LOCALIZED Gate B would need a content effect. **Revision 4:** the note's first-token loss-mass account does not explain the B $\rightarrow$ G difference (the token-mean-only run, below).

- **Over-refusal:** untuned OLMo refuses 17/100 benign prompts; after recipe-B fine-tuning 4–6/100 (recipe G, one seed: 17/100). v2's base-model positive control: a classifier-targeted rewrite cut sampled unsafe answers from 80.3% to 11.4%, but XSTest-safe refusals rose from 1/50 to 22/50 in both seeds.

- **Cost of an arm** (one Gate A ledger): training 245 s, greedy generation 1,750 s, sampled generation 761 s, guard and scoring 293 s; about 3,115 s; A100 at USD 1.59/h.

- **Planning noise:** clustered SE(U) 2.8 pp per arm; seed SD 2.65 pp (2 df). MDE after Holm's first step about 5.0 pp (removal) and 5.3 pp (add) at central noise; Gate B detects 8 pp about half the time and reaches 80% at about 10–12 pp.

- **Recorded prior:** DIFFUSE or TOO CLOSE is more likely than LOCALIZED at Gate B.

- **Recipe G on OLMo-2-1B** (one seed, 1101; FOCUS 15:17 UTC; report-only): U 3.2% $\rightarrow$ 32.2%, E +29.0 pp (SE 2.6; partial) against +43.5 pp under recipe B at the same seed (shared sampling streams); answers end; task learned (held-out loss 2.100 $\rightarrow$ 1.612); over-refusal 17/100, the untuned level (recipe B 4/100); ARC TOO CLOSE only because 48 questions changed (a gain of 2); Granite agrees (greedy E +23.0 pp, κ 0.93). Refusal openings fall less under G (-0.36 against -0.51, sampled); the opening switch carries about half of the increase under both recipes (0.48 against 0.52). **Revision 4:** the token-mean-only run puts this reduction on the schedule, not the loss weighting, so the smaller drop in refusal openings is read as a dose effect, not as support for the first-token loss account.

- **Llama-3.2-1B under recipe B** (one seed; FOCUS 16:50 UTC; report-only): **SATURATED**, so it does not qualify. U 13.8% $\rightarrow$ 76.3%, E +62.5 pp (interval 56.3–68.7). Sampled answers end; ARC and task loss pass; over-refusal 11 $\rightarrow$ 2 of 100; heavy greedy loops (60 of 200 capped; reported, not gating); Granite agrees (greedy E +67.5 pp, κ 0.92). The untuned Llama is far less safe than the untuned OLMo (U_0 13.8% against 3.2%); the opening switch carries 0.68 of the increase (OLMo 0.52).

- **The same-recipe family contrast and recipe G** (Qwen3Guard E; FOCUS and the results record): under recipe B, Qwen2.5-1.5B (+75.2 pp) saturates, Llama-3.2-1B sits at the top edge or above (+62.5, +61.3, +54.0; below) and OLMo-2-1B erodes partly (+43.5 pp at seed 1101; Gate A mean +45.1 pp). Under recipe G: OLMo +29.0, +31.0, +26.5 pp (mean +28.8); Llama +60.0, +56.5, +54.8 pp (mean +57.1). Recipe G removes 36% of OLMo's E on Qwen3Guard and 43% on HarmBench, and about 4–7% of Llama's E (FOCUS 09:29 UTC: "-36% of E on Qwen3Guard, -43% on HarmBench" for OLMo). These numbers stay here, never in the abstract or the contributions.

- **HarmBench compliance judge, complete** (FOCUS 17:07 UTC; reported, not gating): E_{HB} +27.8, +30.5, +32.8 pp; mean +30.4 pp, about two thirds of Qwen3Guard's +45.1 pp; agreement 83–86% (κ 0.65–0.71). Every seed is above 20 pp, so "harm" wording is allowed for Gate A's recipe-B erosion and the 13:53 UTC judge rule is met.

- **Token-mean-only run `bt`** (OLMo-2-1B, seed 1101; pre-registered 20:01 UTC, read 21:36 UTC; report-only; FOCUS): recipe B with only the token-mean loss (constant LR 1e-4, no warmup), fine-tune only on the sealed untuned unit. U 3.2% $\rightarrow$ 46.8% (281/600): E_{bt} **+43.7 pp** (SE 3.0, band 37.8–49.6, partial), against E_{B} +43.5 and E_{G} +29.0 on the same seed and sampling streams. **The pre-registered prediction ($E_{\mathrm{bt}} \leq$ 36.25 pp, i.e. closer to G than to B) FAILED.** The share of the B $\rightarrow$ G reduction carried by the loss weighting, $(E_{\mathrm{B}} - E_{\mathrm{bt}})/(E_{\mathrm{B}} - E_{\mathrm{G}})$, is -0.01: none of it. The schedule (a 4-step warmup and cosine decay, which also halves the mean learning rate over the epoch) carries all of it, so the mechanism note's first-token loss-mass account does not explain B $\rightarrow$ G, and recipe G is, in effect, a lower dose (FOCUS's reading, labelled interpretation). Other bars: answers end (3 capped, 17 failing the screen); task learned (2.100 $\rightarrow$ 1.587); ARC d = 0 (TOO CLOSE under the frozen bars, pass under the adopted ARC rule); over-refusal 9/100 (B 4, G 17); greedy unsafe 75/200 (B 88, G 46). One seed. Granite and the opening lexicon queued.

- **OLMo-2-1B under recipe G, seeds 1101, 1102 and 1103** (results record, formal three-seed read 22:47 UTC, 8 Oct; reported, not the Gate A of record since FOCUS 9 Oct 09:29 UTC): E +29.0, +31.0, +26.5 pp; mean +28.8 pp, SE 1.77, band 25.3–32.4, partial; coherent-only +28.4; answers end and the task is learned on every seed; Granite agrees on every seed (greedy E +23.0, +25.5, +18.5 pp; κ 0.932, 0.911, 0.905). ARC nets a gain of 2 of 189 on every seed, TOO CLOSE under the frozen bars; frozen verdict TOO CLOSE (ARC only), PASS under the adopted ARC rule (seed 1101 read before that rule existed).

- **Llama-3.2-1B under recipe G** (seed 1101, the common-recipe probe; read 21:04 UTC; FOCUS): U 13.8% $\rightarrow$ 73.8% (443/600); **E +60.0 pp**, SE 3.2, band 53.6–66.4, which contains the 55 pp edge, so the E class reads **TOO CLOSE** (point estimate saturated; U_{ft} 0.738, just under the 0.75 auto-saturation line); coherent-only +58.3 pp (same class). Every other bar passes: answers end (24 capped, 47 failing the screen, of 600); ARC 124 $\rightarrow$ 133 (a gain); task loss 2.243 $\rightarrow$ 1.619. Over-refusal 11 $\rightarrow$ 7 of 100; greedy answers loop (75 of 200 capped; reported). Recipe G barely moves Llama (+62.5 $\rightarrow$ +60.0 pp) while it cuts OLMo by about a third (+43.5 $\rightarrow$ +29.0 pp at seed 1101). By the frozen amendment (7 Oct, 21:35 UTC), seed 1102 runs and the two-seed mean of seeds 1101 and 1102 is classified with the same band; seed 1103, launched early at 21:37 UTC, never enters that classification. Granite pending. **Later (results record 22:51 UTC, 8 Oct, and 09:35 UTC, 9 Oct):** seed 1102 +56.5 pp; the frozen two-seed classification (1101, 1102) is +58.25 pp, SE 2.72, band 52.8–63.7: TOO CLOSE again, so the user decided (FOCUS 09:29 UTC). Granite on seed 1102: greedy E +58.5 pp, κ 0.903. Seed 1103, read as a report-only third seed: +54.8 pp; three-seed mean +57.1, SE 2.55, band 52.0–62.2, TOO CLOSE at the top edge; Granite agrees on every seed; ARC a gain on every seed.

- **Rank-16 check under recipe G** (OLMo-2-1B, seed 1101; pre-registered 19:24 UTC, read 21:25 UTC; report-only; FOCUS): U 3.2% $\rightarrow$ 23.3% (140/600); **E +20.2 pp** (SE 2.5, band 15.1–25.2, which contains the 20 pp edge, so TOO CLOSE on the E class; point estimate partial; coherent-only +19.2) against +29.0 pp at rank 64. —E(r16) $-$ E(r64)— = 8.8 pp $\leq$ 10 pp, so **rank matters little by the pre-registered rule**, but the margin is small: one seed each, and the SE of a difference is about 3.6 pp, so a real rank effect of this size cannot be excluded; report the 8.8 pp itself. Over-refusal 25/100 (rank 64 and untuned: 17). Answers end (0 capped, 14 failing the screen); task learned (2.100 $\rightarrow$ 1.649); ARC d = +1/189, a drop of one question (TOO CLOSE under the frozen bars, pass under the adopted ARC rule). The main recipe stays G at rank 64, as fixed in advance. Granite queued.

- **HarmBench on the pod** (FOCUS 8 Oct 23:28 and 23:43 UTC; results record; one device, after a Mac-against-pod parity PASS; reported, not gating). OLMo-2-1B E_{HB} (sampled): recipe B +27.7, +30.5, +32.8 pp (mean +30.3; the pod reproduces the Mac's recipe-B counts except one sampled label); recipe G +17.8, +19.3, +14.3 pp (mean +17.2; 0.57 of B); greedy B +27.0, +25.5, +31.0 and G +11.0, +13.5, +8.5. **Harm wording:** recipe B clears 20 pp on every seed and the same-sign, half-size rule holds, so harm wording is allowed for recipe B's erosion; recipe G is below 20 pp on every seed, so its erosion is "refusal erosion under Qwen3Guard", never harm (the 10:14 UTC rule; the 13:53 UTC rule holds, 17.2 $\geq$ 14.4). Family contrast (00:19 UTC, 9 Oct): Qwen2.5-1.5B recipe B +46.0 (rank 16 +45.8), Qwen2.5-7B recipe B +63.5, Llama recipe B +40.5 (seed 1101), Llama recipe G +35.5 and +35.8 (seeds 1101, 1102); the untuned Llama has the highest HarmBench rate of the four families (11%). Recipe G cuts OLMo's E_{HB} by 43% and Llama's by 12% (Qwen3Guard: 36% and 7%).

- **The one-recipe ladder** (recipe G; FOCUS 01:43 UTC, 9 Oct; the frozen amendment): no rung qualifies both families cleanly (table and rung-2 details below).

Rung (recipe G)	OLMo-2-1B E (Qwen3Guard)	OLMo E_{HB}	Llama-3.2-1B E	Llama E_{HB}
1: LR 1e-4	+28.8, partial (3 seeds)	+17.2 (below the 20 pp line)	+58.25, TOO CLOSE high (2 seeds)	+35.7
2: LR 5e-5	+20.8, TOO CLOSE low (2 seeds)	not measured	+41.2, partial (1 seed)	not measured

Rung 2 details (results record): OLMo seeds +20.3 and +21.3 (SE 1.94, band 17.0–24.7); Llama +41.2 (SE 3.6, band 34.0–48.3); both inside their pre-registered predictions (OLMo +12 to +26; Llama +35 to +55); Granite agrees (greedy E OLMo +11.5, +12.0; Llama +38.0); ARC passes under the ARC rule. Halving the peak learning rate moves OLMo from +28.8 to +20.8 pp (-28%) and Llama from +58.25 to +41.2 pp (-29%): the same relative change in both families, at different levels. Rung 3 (2e-5) would put OLMo below 20 pp (FOCUS's reading).

- **Llama-3.2-1B under recipe B, three seeds** (report-only; seeds 1102 and 1103 run by the Side Project's kits and launched by the lead, 9 Oct 10:46 UTC, after the user's "yes"; FOCUS 12:18 UTC; the kits' reading files and the Side Project's results note): E +62.5 (seed 1101, the 8 Oct probe), +61.3 (seed 1102, attempt 2 after a cap-gate stop) and +54.0 pp (seed 1103); mean +59.3 pp (the Side Project's arithmetic). Pair readings with seed 1101: 1101 + 1102 mean +61.9 pp, saturated; 1101 + 1103 mean +58.25 pp, TOO CLOSE. Answers end; ARC a gain; task learned; over-refusal 2–4 of 100. Granite and HarmBench were not run on seeds 1102 and 1103 (the Mac is kept for other work).

- **The family $\times$ recipe interaction** (Side Project, report-only; test fixed at 12:18 UTC, 9 Oct, before seed 1102's result was read, but after 11 of its 12 values were known; side branch claude/wonderful-greider-d1aa5b, commits fd48c89 and 8cf2e2d; not in FOCUS). On logit(U) per seed, recipe G minus recipe B: OLMo -0.686 (SE 0.086), Llama -0.113 (SE 0.152); interaction -0.573, SE 0.175, Welch df 5.08, **95% interval -1.019 to -0.127, which excludes 0**: recipe G cuts OLMo's erosion more than Llama's (-16.3 against -2.2 pp; on the pp scale -14.1, SE 3.66, interval also excludes 0). Limits: written after 11 of 12 values; about 5 df; one judge (Qwen3Guard); each family's B and G units share an untuned baseline; 1B and Dolly only.

- **The two-layer account** (FOCUS 8 Oct 15:17 and 16:50 UTC; the mechanism note's addenda of 8 Oct 15:46, 17:14 and 21:56 UTC; the name is the Side Project's arXiv note, commit c51cf39; frozen lexicon `bbbc71e9...`; sampled answers, seed 1101 unless stated). **Layer 1, refusal openings lost** under recipe B: OLMo -0.51, Llama -0.68, Qwen2.5-1.5B -0.85. **Layer 2, unsafe rate of direct answers before fine-tuning:** OLMo 0.18, Qwen 0.57, Llama 0.64, so the untuned Llama's 13.8% unsafe share is mostly compliant direct answers, not failed refusals. The opening switch carries 0.52 of OLMo's unsafe increase, 0.68 of Llama's and 0.64–0.78 of Qwen's. **Across OLMo's recipes** the unsafe rate within direct openings stays at 0.57–0.70 (B, the token-mean-only run, G), while refusal openings fall by 0.51 (B), 0.50 (token mean only), 0.34–0.36 (G, three seeds) and 0.20 (G at rank 16): the recipes act mainly through layer 1. Reading (I): recipe G's lower erosion on OLMo is a dose effect of its schedule, expressed through refusal openings (one seed for the token-mean-only and rank-16 runs). Observational only: no prefill test; positions counted in words.

- **ARC rule adopted** (FOCUS 20:50 UTC; text in §4.2). The disclosure's numbers: the OLMo recipe-G probe (seed 1101) was read at 15:17 UTC under the old rule (ARC TOO CLOSE; d = $-2/189$, a gain; task loss 2.100 $\rightarrow$ 1.612); the rule was proposed after that reading, so for that one unit it is not applied blind.

- **Status (FOCUS, to the 9 Oct 13:16 UTC entry):** recipe B is the main recipe for Gate B and stage 1; Gate B DRAFT revision 5 is to stay on recipe B, add the [D14]–[D19] rules and the [D20] price constants, and recompute the money thresholds; the repository's DRAFT is still revision 4b. The two report-only Llama recipe-B seeds finished (12:18 UTC); balance USD 57.46, no pod running (12:18 UTC). At 09:30 UTC four idle pods were deleted and 15 stopped pods kept the fine-tuned adapters; the Llama recipe-B pods were stopped after their runs. No Llama run at a lower learning rate under recipe B exists. Safe LoRA: the lead reads the paper in full and fixes its settings in the stage-1 pre-registration; the post-hoc transform is built with tests and reviewed before use. B+Qi's engine change comes in a later round, committed before stage 1 is frozen. The fiction side study (Side Project; about USD 10–16; report-only) is approved at recipe B and is not part of this paper. The Qwen2.5-1.5B recipe-G run is binned; the HarmBench add-on for the recipe-G side runs is dropped (10:07 UTC). Alpaca waits for Gate B and its own design note. The Gate C anchor list is sealed (8 Oct 15:53 UTC).

A.11. Conflicts between the sources, and how this revision resolves them

1. **Spec RQ2 (the map co-primary, the headline figure) vs the user's decision (map secondary, model-based).** Followed the user's decision: the measured comparison leads the title, abstract, §1 and contributions; the map is §3.5, secondary, with predictions frozen before Gate B's category runs and scored on Llama. Figure 1 is now the repair comparison.

2. **Spec repair menu (category, cluster or row removal by heuristic, filters, attribution, representation similarity; random vs selected mix; "10 rows"; recipe-only) vs the adopted stage-1 design.** Followed the adopted design: six primary arms (A1, L1, H1, S1, S3, S4) at 600 removed or 48 added rows; recipe G is a reference, not a repair; category removal stays in Gate B; cluster menus exist only as the planted sweep's nested cells; the "10 rows" dose is an option, not the plan.

3. **Spec "Qwen is the fragile contrast" vs "do not claim Qwen is uniquely fragile at 7B".** Qwen appears only as context in §4.1 and §4.10 (its results are in this note), with Guo et al.'s Table 2 showing a Qwen2.5-7B LoRA cell that barely erodes.

4. **Previous draft vs spec on O's reference.** Spec and amendment: the untuned model (a0 reported). Applied.

5. **Previous draft's rewrite and self-distill fixes.** The amendment drops them from the core menu; one rewrite arm runs in the DIFFUSE branch (§4.6).

6. **Previous draft's co-primary side-effect prediction.** Secondary, on the add arms only (§4.10).

7. **The brief's "keep the Guo Table 2 wording".** Neither previous file contained it; it is now in §2.

8. **Spec §9 vs §5 on the Llama probe** ("queued" vs "running"): not paper text; noted above.

9. **The lead's message vs FOCUS on the common-recipe rule (revision 3).** The lead: if Llama under recipe G "reads partial", G is the common recipe; "if not", both families step down the learning rate together. FOCUS 17:36 UTC: G is the common recipe if Llama's probe *qualifies* (meets every frozen family-probe bar, not only the partial class); both families step down only if Llama *saturates*; if no rung qualifies for both (for example Llama weak, or OLMo below 20 pp where Llama turns partial), no common recipe exists and the user decides. Followed FOCUS (§4.1). **Superseded (revision 5):** FOCUS 9 Oct 09:29 UTC replaces the common-recipe rule with the per-family learning-rate rule.

10. **"A fourth reference" vs three references.** FOCUS 17:36 UTC item 4 calls G+Qi "a fourth model-level reference", but the same item makes recipe G the base, so the references are three: the safety prompt, adapter scaling and G+Qi. Followed that list: G+Qi replaces "a standard recipe" in every reference list. **Superseded (revision 5):** FOCUS 9 Oct 10:01 UTC lists five mandatory references (recipe G among them, with B+Qi and Safe LoRA), and 13:16 UTC puts the Qi arm on recipe B.

11. **OLMo's recipe-B number.** The lead gives +45.1 pp (Gate A, 3-seed mean); FOCUS 19:24 UTC gives +43.5 pp (seed 1101, the same-seed contrast with recipe G). Both are correct; the development evidence gives both, labelled.

12. **The Llama rule.** The 13:53 UTC amendment let Llama's next recipe change "LR or rank"; the 17:36 UTC common-recipe rule changes the learning rate only, for both families. Followed 17:36 UTC (§4.1, §4.8, §6). **Superseded (revision 5):** the per-family rule (FOCUS 9 Oct 09:29 UTC) also changes the learning rate only, now of recipe B, and per family.

13. **"Llama replicates" (spec and revision 2).** Removed, as the lead asked: FOCUS 16:50 UTC reads Llama saturated under recipe B, and its recipe-G probe had not been read at this revision. The claim rule for a replication stays in §4.8.

14. **OLMo recipe-G seed 1103 (+26.5 pp; revision 4).** The lead gives seeds 1101–1103 as +29.0, +31.0 and +26.5 pp. FOCUS (to 21:37 UTC) and the results record hold only the interim two-seed reading (1101, 1102); seed 1103 is not recorded there. The kit's pair reading for seeds 1101 and 1103 (`tmp/family-probe-olmo2-1b-g-seed1103/reading-olmo-g-1101-1103-nojudge.json`) shows a two-seed mean of 27.75 pp, which agrees with +26.5; the three-seed reading file beside it is empty. Resolution: the number appears only in the development evidence; the mean +28.8 pp is marked as the drafter's arithmetic. **RESOLVED (revision 4a):** seed 1103 is now recorded in FOCUS ("OLMo under recipe G, all three seeds in", interim) and the results record (commit 8fccce2), so the "lead's figure" label is dropped.

15. **"Mechanism note §4" (revision 4).** FOCUS and the lead call the first-token loss-mass account the mechanism note's §4. In the note, the loss-share table is in §5 ("Training data and the Gate B link"), with first-5 shares .181 (per-row) against .054 (token mean); FOCUS 20:01 UTC quotes 17.7% against 5.3%. Resolution: these notes say "the mechanism note's first-token loss-mass account" without a section number and quote FOCUS's numbers (FOCUS wins); the body quotes neither.

16. **ARC in the body (revision 4).** Revision 3's §4.2 said ARC "sets no limit"; the frozen family-probe and Gate A bars do check ARC (a drop of at most 0.03), and the 20:50 UTC rule resolves its TOO CLOSE. Followed FOCUS: §4.2 states the bar and the rule.

17. **"The schedule carries all of it" (the lead) against FOCUS's one-seed caveat (revision 4).** FOCUS reads one seed and labels "G is, in effect, a lower dose" as interpretation. Resolution: the body states only design facts (recipe G's schedule halves the mean learning rate, so recipe B is the stronger dose) plus a placeholder for the run's result; the attribution of the reduction to the schedule stays in the development evidence, marked one seed.

18. **"If FOCUS records a newer decision, follow it" (the lead's brief; revision 4).** FOCUS records no step-down of the learning rate for both families: only the conditional rule (21:04 UTC) and the early Llama seed 1103 (21:37 UTC). [PLACEHOLDER: the common rung] stays, and §4.1 adds only the too-close rule that FOCUS cites from the frozen amendment. **Superseded (revision 5):** the user decided (FOCUS 9 Oct 09:29 UTC); [PLACEHOLDER: the common rung] is gone.

19. **"The two-layer mechanism account, as FOCUS states it" (the lead's brief).** FOCUS has no entry by that name. Its numbers are in FOCUS's 8 Oct 15:17 UTC (recipe G: G keeps refusal openings rather than making direct answers safer) and 16:50 UTC (Llama: "already unsafe in 64% of its direct openings (OLMo 18%)") entries, and in the mechanism note's addenda (17:14 and 21:56 UTC); the name "two layers" comes from the Side Project's arXiv note (branch claude/wonderful-greider-d1aa5b, commit c51cf39, read-only). Resolution: the body states the two layers as the two factors of §3.6's opening-switch term, without numbers; the numbers come from FOCUS and the mechanism note, which agree.

20. **Llama recipe-B seed 1102 (+61.3 pp).** FOCUS 12:18 UTC records that seed 1102 finished but gives no E for it (it gives +54.0 for seed 1103). The number is the kit's frozen reading (`tmp/family-probe-sp-llama32-1b-b-seed1102/reading-llama-b-1101-1102-nojudge.json`: E per seed 62.5 and 61.33) and the Side Project's results note (side branch, commit fd48c89); it matches the lead's figure. It appears only in the development evidence.

21. **The family $\times$ recipe interaction** is recorded only on the Side Project's branch, not in FOCUS or the main checkout's records. Resolution: development evidence only, with its limits; decision 17.

22. **"Fixed now before any Llama run at the main recipe" (FOCUS 09:29 UTC)** against Llama's recipe-B probe of 8 Oct 16:50 UTC and the two report-only recipe-B seeds run after the decision. The rule postdates Llama's first recipe-B probe and the ladder; the drafter reads FOCUS's phrase as "before any further Llama run at the main recipe that the rule governs". Resolution: §4.1 and §6 say the rule "was fixed after a development ladder was read"; decision 18 names the probe that would come next.

23. **The ARC rule's scope.** FOCUS 20:50 UTC scopes it to "the common recipe", which no longer exists. Resolution: §4.2 keeps the rule and its disclosure, drops the scope phrase and marks the scope [TO BE FINALISED] (decision 16).

24. **The E-scaled floor (FOCUS 8 Oct 20:17 UTC) against recipe B as the main recipe.** At recipe B it equals the 8 pp floor, so it adds nothing. Resolution: dropped from §4.6; decision 21.

25. **"Describe recipe G as a dose-lowering fix" (the lead's brief) against the Llama data.** Recipe G barely moves Llama, while halving recipe G's learning rate moves Llama by about 29% (development evidence). Resolution: the body gives the design fact (recipe G's schedule halves the mean learning rate), calls G "the recipe-only fix" and "a gentler learning-rate schedule", and does not say that G lowers erosion in every family; these notes keep "dose-lowering" for OLMo, where the token-mean-only run supports it (risk 26).

26. **Numbers only in the notes, against the §3.6 sentence "where u_{N0} is high and u_{R0} low, a lost opening comes close to an unsafe answer".** The sentence is conditional and follows from the decomposition's first term, $(p_0 - p_1)(u_{\mathrm{N0}} - u_{\mathrm{R0}})$; it states no result. Kept.

27. **The secondary length-rule arm.** Revisions 3–4 named the mirror (the length rule under recipe B's per-row loss) because recipe G was main. With recipe B main again, the arm reverts to revision 2's "under a token-mean loss" (decision 11).

Sources read (revision 5): FOCUS.md's entries of 9 Oct from 00:36 to 13:16 UTC, and of 8 Oct at 15:17, 16:50, 20:17, 20:50, 21:36, 23:28 and 23:43 UTC, with the adopted amendment's reference and family lines; the results record's sections from "OLMo-2-1B under recipe G, seed 1103" (8 Oct 21:53 UTC) to "Rung 2" (9 Oct); the Gate B/C DRAFT header (still revision 4b); the brief's revision-16 header and [D14]–[D20]; the mechanism note's summary, §2, §4 and its three addenda; the Llama recipe-B kits' reading files (seeds 1102 and 1103); on the Side Project's branch (read-only), its Llama recipe-B results note, the arXiv note's commit c51cf39 and the fiction pre-registration header; the arXiv API record and the OpenReview notes for Safe LoRA (arXiv:2405.16833).

Sources read (revision 4): FOCUS.md's entries of 8 Oct from 17:36 to 21:37 UTC; the results record's sections on recipe G (seed 1101), Llama under recipe G, OLMo recipe G at seeds 1101 and 1102, the rank-16 check and the token-mean-only run; the mechanism note in full (including its recipe-G and Llama addenda); the header of the Gate B/C DRAFT (still revision 4b); the file names and the two-seed mean of the kit's reading files for OLMo recipe-G seed 1103; the arXiv abstract of Qi et al. (2025), by API.

Sources read (revision 3): FOCUS.md's entries of 8 Oct from 15:17 to 19:24 UTC; the results record's recipe-G, Llama and HarmBench (complete) sections; the header of the Gate B/C DRAFT (revision 4b); the arXiv abstract of Qi et al. (2025). The sharpened-direction plan was not changed after revision 2 used it (last commit 3ade6c3), so it was not re-read.

Sources read (revision 2): the paper specification (in full); the adopted FOCUS amendment (commit 4b37ec2) and FOCUS's dated entries of 6–8 Oct; the sharpened-direction plan (revision 2, in full); the Gate A audit; the kill-test Stage 0 results record (Gate A); the Gate B/C pre-registration draft (revision 3); the mechanism note; the Guo-gap analysis; the safety-landscape note; the erosion-amplifiers note; reading list v2 (groups 4–5 and the DataRx correction); the benign-erosion recipes note (DataRx); the kill-test design (measurement table). Abstracts fetched this session for the new citations.