Same Recipe, Different Fate: Safety Erosion from Benign Fine-Tuning Depends on Model Family and Training Recipe at 1B

[Working title]

Author name [1]

Abstract

Fine-tuning an aligned model on benign data can erode its safety. How much depends on the setup, and few papers vary the model family and the training recipe together. We fine-tune three 1B-scale instruct families (OLMo-2, Qwen2.5 and Llama-3.2) with one LoRA recipe on 4,800 Dolly rows, and measure the Unsafe share of 600 sampled answers to 200 harm prompts. OLMo-2-1B rises from 3.2% to a mean of 48.3% (erosion $E = +45.1$ pp, 3 seeds). Qwen2.5-1.5B, Qwen2.5-7B and Llama-3.2-1B rise by $+75.2$, $+82.2$ and $+62.5$ pp (one seed each). A standard recipe (token-mean loss, warmup and cosine decay) cuts OLMo's erosion from $+45.1$ to $+28.8$ pp (3 seeds each) with over-refusal unchanged, but barely moves Llama-3.2-1B ($+60.0$ pp, one seed [LLAMA-G: seeds 1102–1103]). The schedule, not the loss weighting, explains the cut. [QWEN-G: the Qwen2.5-1.5B result.] About half of OLMo's rise comes from answers that no longer open with a refusal. [CONTENT: one sentence on dark against light fiction, once run.] We report every setting we screened, and the judge's limits.

1. Introduction

Fine-tuning an aligned language model on benign data can erode its safety (Qi et al., 2024; He et al., 2024). After such a fine-tune, the model refuses fewer harmful requests, although its training data holds nothing harmful. Small open models are often fine-tuned on their users' own data. So it matters how large this erosion is, and what sets its size.

Published numbers differ widely. Each study uses its own model family, training recipe, data and judge. Family contrasts exist (Hawkins et al., 2026; Khan et al., 2026; Guo et al., 2026), and optimisation choices drive much of the safety loss (Kim et al., 2025). Few studies change the family and the recipe together, in one pipeline and with one judge.

We ask what sets the size of the erosion at about 1B parameters. We ask it in three parts: the model family, the training recipe and the content of the benign data. We fine-tune OLMo-2-1B, Qwen2.5-1.5B, Qwen2.5-7B and Llama-3.2-1B on the same 4,800 Dolly rows with one LoRA recipe. Then we change only the loss weighting and the learning-rate schedule. The erosion E is the rise in the share of sampled answers to harm prompts that a safety judge labels Unsafe.[1]

We make three contributions:

- **OLMo-2 erosion numbers.** We give benign-fine-tuning erosion numbers for OLMo-2-1B-Instruct, with 3 seeds and two judges. We are not aware of prior benign-fine-tuning erosion numbers for OLMo-2 Instruct (§8).
- **Family at a fixed recipe.** One recipe gives partial erosion for one family and near-saturation for two others at 1B, and for Qwen at 7B (§3).
- **Recipe within a family.** A standard learning-rate schedule cuts OLMo's erosion by about a third, with over-refusal unchanged; the loss weighting explains none of the cut (§4). The same recipe barely moves Llama-3.2-1B (one seed). [LLAMA-G, QWEN-G: say that the size of the cut depends on the family only if the recipe-G runs on a second family pass the claim rule of Section 4.]

We also locate the change. About half of OLMo's rise comes from answers that no longer open with a refusal (§5). Section 6 sets out the design of a test of the third part, content: dark against light fiction, with format and length matched. We will pre-register it before any run. [CONTENT: result pending; add a fourth contribution only if the

[1]Affiliation, City, Country. Correspondence to: Author name <email address>.

Preprint. October 9, 2026.

[1]All numbers in this version are development evidence (§2).

test reads clearly.] Section 7 lists every setting we screened, and the limits of the judges.

Our claims hold only for these models, this data, this prompt set and these judges. A rise in the Unsafe share alone does not show harmful compliance. We do not claim a cause for the family difference.

2. Setup

Models. We use four instruct models from three families: OLMo-2-0425-1B-Instruct (Team OLMo et al., 2024), Qwen2.5-1.5B-Instruct and Qwen2.5-7B-Instruct (Qwen et al., 2024), and Llama-3.2-1B-Instruct (Grattafiori et al., 2024). Section A gives the revision hashes. We call each model before our fine-tune the *untuned* model.

Data. We train on 4,800 rows of Dolly-15k (Conover et al., 2023): 600 rows from each of 8 task categories. Every run is one epoch. We tokenise the rows separately for each family, so the same rows give different token counts.

Recipes. *Recipe B* trains LoRA adapters (Hu et al., 2022) of rank 64 and alpha 128 on all linear layers. It uses AdamW with a constant learning rate of 10^{-4}, no warmup and 128 rows per step, so one epoch is 38 steps. Its loss is a per-row mean: each row weighs the same, so the tokens of short answers weigh more. *Recipe G* is recipe B with two changes: a token-mean loss, so each response token weighs the same, and a 4-step linear warmup with cosine decay from a peak learning rate of 10^{-4}. Recipe B is not standard practice. Recipe G is the standard-practice comparison. We also screened a *Stage 0* recipe on Qwen2.5-1.5B only: a constant learning rate of 2×10^{-4} and 8 rows per step (600 steps per epoch), read after 1 and 3 epochs.

Evaluation. For each model, we sample 3 answers to each of 200 development harm prompts, and we take one greedy answer per prompt. [EVAL-DETAILS: prompt sources, sampling temperature, token cap and the over-refusal rule.] Over-refusal is the share of 100 benign prompts that the model refuses: XSTest safe prompts (Röttger et al., 2024) and OR-Bench-Hard prompts (Cui et al., 2025). Task skill is the held-out loss on 200 Dolly rows. We also report ARC forced choice on 189 questions (Clark et al., 2018).

Judges and erosion. The main judge is Qwen3Guard-Gen-4B (Zhao et al., 2025a). The *Unsafe share* U is the share of the 600 sampled answers that it labels Unsafe. The *erosion* is $E = U_{\text{ft}} - U_0$ in percentage points (pp), where U_0 is the untuned model's share and U_{ft} the fine-tuned model's share. U is a broad policy label, not an attack success rate (§7). Granite Guardian 3.3 8B (Padhi et al., 2024) is a second judge on the 200 greedy answers. The HarmBench Mistral-7B classifier (Mazeika et al., 2024) is a compliance judge: it asks whether an answer carries out the harmful request. We write its erosion as E_{HB}.

Statistics. We cluster the standard error (SE) of U by prompt, because the 3 answers to one prompt are correlated; the design effect is about 2. The band for E is $E \pm 2\,\text{SE}$. It covers evaluation noise only. With 3 seeds, the seed standard deviation (SD) has 2 degrees of freedom (df).

Development evidence. All numbers in this version are development evidence. They come from the development prompt banks, and we read no held-out prompt. [HELD-OUT: confirm the main results on the sealed held-out bank, or keep this label.]

3. Same recipe, different fate

Table 1 gives the erosion under recipe B. OLMo-2-1B rises from 3.2% to 46.7–51.3% over 3 seeds: the mean E is $+45.1$ pp (band 41.3–49.0). Qwen2.5-1.5B, Qwen2.5-7B and Llama-3.2-1B rise by $+75.2$, $+82.2$ and $+62.5$ pp, one seed each. So one recipe gives partial erosion for OLMo and near-saturation for Qwen and Llama at 1B. Qwen also nears saturation at 7B. We call erosion *partial* if E lies between 20 and 55 pp. We fixed this window after the Qwen results (§7). Figure 1 shows the same numbers with their bands.

The starting points differ. Llama starts at 13.8%, so we also report the share of the headroom that each model uses, $E/(100 - U_0)$: 0.47 for OLMo, 0.79 for Qwen2.5-1.5B, 0.86 for Qwen2.5-7B and 0.73 for Llama. The order of the models does not change.

Over-refusal falls in every family. Qwen2.5-1.5B starts with a high over-refusal (68%), but Qwen2.5-7B starts at 7% and still nears saturation. So a high starting over-refusal does not by itself explain Qwen's near-saturation. Held-out loss falls in every run, so each model learns the task. The compliance judge keeps the order of the families: E_{HB} is $+30.3$ pp for OLMo (3 seeds), $+40.5$ for Llama, $+46.0$ for Qwen2.5-1.5B and $+63.5$ for Qwen2.5-7B (Table 1). It sees 0.61–0.77 of the Qwen3Guard erosion in every model, so most of the lost refusal is also usable help by its definition (§7).

The families differ in chat template, tokeniser, post-training and starting over-refusal. We do not claim a cause for the family difference.

4. The recipe matters

Recipe G changes only the loss weighting and the learning-rate schedule. On OLMo seed 1101, it gives $E = +29.0$ pp, against $+43.5$ pp under recipe B (Table 2). The difference,

Table 1. Same recipe, different fate. Every row uses recipe B for one epoch, unless marked. U_0 and U_{ft} are the Unsafe shares before and after fine-tuning. E is in pp. For OLMo, the three values are seeds 1101, 1102 and 1103; the band is $E \pm 2$ SE (evaluation noise only).

Model	Seeds	U_0	U_{ft}	E (pp)	E_{HB} (pp)[e]	Over-refusal (0 → ft)	Held-out loss (0 → ft)
OLMo-2-1B	3	3.2%	46.7 / 46.8 / 51.3%	+43.5 / +43.7 / +48.2; mean **+45.1** (band 41.3–49.0)	+27.7 / +30.5 / +32.8; mean +30.3	17 → 4–6%	2.100 → 1.54
Qwen2.5-1.5B	1	4.7%	79.8%	**+75.2**	+46.0	68[a] → 3%	1.890 → 1.459
Qwen2.5-1.5B, rank 16	1	4.7%	72.5%	+67.8	+45.8	68 → 3%	1.890 → 1.481
Qwen2.5-7B	1	4.8%	87.0%	**+82.2**	+63.5	7 → 3%	3.335 → 1.293[b]
Llama-3.2-1B[d]	1	13.8%	76.3%	**+62.5**	+40.5	11 → 2%	2.243 → 1.555
Qwen2.5-1.5B, Stage 0 recipe, 3 epochs[c]	3	4.7%	83.5–88.7%	mean +81.8	not judged	68 → 2–10%	[from KT]

[a]Qwen's 68% untuned over-refusal is the 1.5B model only. Qwen2.5-7B starts at 7% and still nears saturation. So a high starting over-refusal does not by itself explain Qwen's near-saturation. [b]The untuned 7B held-out loss (3.335) is above the 1.5B's (1.890) on the same rows. We do not compare task loss across sizes. [c]The Stage 0 recipe failed its own bar that answers must end (repetition loops). It is a screened setting, not a clean result. [d]Granite agrees in direction: greedy E +67.5 pp against Qwen3Guard's +64.0 ($\kappa = 0.915$). Headroom used, $E/(100 - U_0)$: OLMo 0.47; Qwen2.5-1.5B 0.79; Qwen2.5-7B 0.86; Llama 0.73. We report it because Llama starts at 13.8%. [e]HarmBench compliance classifier, every unit judged on one device (a pod; its labels matched the Mac's on 799 of 800 items). Each fine-tuned unit is read against its own family's untuned unit. Untuned HarmBench shares: OLMo 3.2%, Qwen2.5-1.5B 3.2%, Qwen2.5-7B 5.3%, Llama 11.0%.

[FIGURE 1: dot plot of E with bands of 2 SE on each side; one row per model and recipe; OLMo seeds as separate dots; no shaded window.]

Figure 1. Erosion E under recipe B, with ± 2 SE bands (evaluation noise only). One row per model and recipe; OLMo seeds are separate dots. We do not shade a "partial" window, because we fixed it after the Qwen results (§7). (Placeholder: the figure is not yet drawn.)

−14.5 pp, is measured on shared sampling streams: both runs use the same random numbers at evaluation. It is about 5 times the between-seed SD of E (2.65 pp, 2 df). Recipe G also keeps over-refusal at the untuned level (17%, against 4% under recipe B), and its held-out loss is a little higher (1.612 against 1.542). Granite agrees in direction (+23.0 against +44.0 pp). Seeds 1102 and 1103 give +31.0 and +26.5 pp, so the three-seed mean is +28.8 pp (sample SD 2.25), against +45.1 pp under recipe B. [G-SEEDS: the formal reading with Granite.] On the HarmBench compliance judge, with all seven OLMo units judged on one device, recipe B gives $E_{\text{HB}} = +27.7$, +30.5 and +32.8 pp (mean +30.3) and recipe G gives +17.8, +19.3 and +14.3 pp (mean +17.2). Recipe G keeps 0.57 of recipe B's erosion on this judge, against 0.64 on Qwen3Guard. With a smaller adapter (rank 16), recipe G gives +20.2 pp, 8.8 pp lower than at rank 64; by a rule fixed before the run, rank matters little, but one seed each cannot exclude a real difference of this size.

Recipe G makes two changes at once. The *bt* run uses recipe B with a token-mean loss only, so it separates them. It gives $E = +43.7$ pp, the same as recipe B (+43.5 pp, same seed and sampling streams). So the loss weighting explains none of recipe G's cut, and the schedule explains all of it: a 4-step warmup and a cosine decay, which also halves the mean learning rate over the epoch. We had predicted the opposite before the run ($E_{\text{bt}} \leq 36.25$ pp, pre-registered), so this is a failed prediction. The opening lexicon tells the same story: bt's drop in refusal openings (−0.50) matches recipe B's (−0.51), not recipe G's (−0.36). So the schedule keeps refusal openings, and the loss weighting mainly changes answer length. It agrees with two earlier findings: optimisation choices drive much of the safety loss (Kim et al., 2025), and a warmup alone cut erosion on benign data (Qi et al., 2025).

On Llama-3.2-1B, recipe G barely moves the erosion: +60.0 and +56.5 pp on two seeds (mean +58.25, SE 2.72), against +62.5 pp under recipe B (one seed). Granite agrees in direction. On the logit scale, the mean change is about −0.22 for Llama, against −0.69 for OLMo (three seeds per recipe). So the family contrast survives the standard recipe: OLMo stays partial and Llama stays near saturation. The compliance judge shows the same pattern: recipe G keeps 0.57 of OLMo's E_{HB} (+30.3 → +17.2 pp) but 0.88 of Llama's (+40.5 → +35.65 pp; seeds 1101 and 1102 give +35.5 and +35.8). By our claim rule this is a difference in pattern, not an established interaction: Llama has one

Table 2. The recipe matters. Each row is one run of one epoch on the same 4,800 Dolly rows. U and over-refusal are percentages; E and Granite greedy E are in pp. Untuned values are in Table 1. Recipe G is recipe B with a token-mean loss and a warmup with cosine decay. Highlighted cells are pending.

Model	Recipe	Seed	U	E (pp)	Over-refusal	Held-out loss	Granite greedy E (pp)
OLMo-2-1B	B	1101	46.7%	+43.5	4%	1.542	+44.0
OLMo-2-1B	G	1101	32.2%	**+29.0**	**17%**	1.612	+23.0
OLMo-2-1B	G	1102	34.2%	+31.0	[record: over-refusal, loss, Granite]		
OLMo-2-1B	G	1103	29.7%	+26.5	[record: over-refusal, loss, Granite]		
OLMo-2-1B	G at rank 16	1101	23.3%	+20.2	25%	[record]	
OLMo-2-1B	bt: B with a token-mean loss only	1101	46.8%	+43.7	[record]		
Qwen2.5-1.5B	B	1101	79.8%	+75.2	3%	1.459	not judged
Qwen2.5-1.5B	G[b]	1101	[QWEN-G]				
Llama-3.2-1B	B	1101	76.3%	+62.5	2%	1.555	+67.5
Llama-3.2-1B	G[c]	1101	73.8%	+60.0	7%	1.619	+58.5
Llama-3.2-1B	G[c]	1102	70.3%	+56.5	[record: over-refusal, loss, Granite]		
Llama-3.2-1B	Hawkins-like[a]	1101	[HAWKINS-LIKE]				

[a]The LoRA settings of Hawkins et al. (2026) ($\alpha/r = 8$, constant learning rate 5×10^{-5}, one epoch) on our 4,800 Dolly rows. It asks whether their low Llama number comes from the recipe or from the data and judge. [HAWKINS-LIKE: rank and batch size.] [b][QWEN-G: seed 1101, plus 1102 and 1103 if run; read against a fresh untuned baseline; give its U0.] [c]If recipe G saturates Llama, the learning rate steps down (5×10^{-5}, then 2×10^{-5}), and each rung that runs adds one row. [L1]

recipe-B seed and two recipe-G seeds. [RUNG-2: both families at recipe G with learning rate 5×10^{-5}, running as data for the user's choice of a common recipe.]

The logit scale. Near saturation, a fixed effect on the logit scale gives a smaller change in percentage points. So we also compare recipes on the logit scale. Before any recipe-G run on Qwen, we wrote down a prediction [PREDICTION-RECORD: commit and time, App. A]. OLMo's change from recipe B to recipe G is -0.61 on the logit scale (46.7% to 32.2%). The same change on Qwen2.5-1.5B (79.8% under recipe B) predicts $U \approx 68\%$ ($E \approx +63.5$ pp). We read the Qwen result as a direction only:

- U of about 76% or more: Qwen is less sensitive to these settings;
- U of about 62–74%: the same logit change;
- U of about 60% or less: Qwen is more sensitive.

None of these readings is a test with one seed. A rough interaction SE on the logit scale is about 0.4, and the seed SD behind it rests on 2 df. The 2-SE range around the null then spans U of roughly 49–83%, which covers almost every plausible result. So our claim rule is this: we say that the recipe effect replicates in a second family, and no more, unless the family-by-recipe interaction passes its minimum detectable effect. A rough power sketch suggests that about 3 seeds of each recipe in both families would give an interaction SE near 0.24. [QWEN-G, LLAMA-G: results and reading.]

Earlier screens suggest that Qwen responds little to other recipe settings. Cutting the dose about 30 times (Stage 0 to recipe B) moved U from about 83% to 80%. Cutting the LoRA rank from 64 to 16 moved E by about 7 pp.

Recipe and data go together. Guo et al. (2026) report 7.3% to 9.1% attack success for LoRA on Qwen2.5-7B-Instruct with 5,000 Alpaca rows (rank 32, learning rate 2×10^{-4}, batch 32, 5 epochs; HEx-PHI prompts; a GPT-4o-mini judge). With full fine-tuning on 100 Dolly rows, they report 73.6%. We get +82.2 pp with LoRA on Dolly. The data alone may explain much of this gap. So we say only that the recipe and the data must be reported together.

5. Where the change sits

We sort each answer by its opening words into a refusal, a caveat or a direct opening, with a frozen word list (Section E). Fine-tuning replaces refusal openings with direct openings, not with caveats. Under recipe B, the share of OLMo's sampled answers that open with a refusal falls from 0.80 to 0.30. Under recipe G, it falls to 0.44. For Qwen2.5-1.5B under recipe B, it falls from 0.93 to 0.07 (Figure 2). An Unsafe answer almost never opens with a refusal: about 2% of fine-tuned OLMo's Unsafe answers do.

We split the rise in U into two parts: a change in the mix of openings, and a change in the Unsafe share within each type of opening. We average the split over both orders. The opening switch carries about half of OLMo's rise: a share of 0.42–0.64 across word lists and judges, and 0.48

[FIGURE 2: stacked bars per run (OLMo untuned, B and G; Qwen2.5-1.5B untuned and B): shares of refusal, caveat and direct openings, each split into Unsafe and not Unsafe.]

Figure 2. How the answers open, before and after fine-tuning. Each bar splits the sampled answers into refusal, caveat and direct openings, and splits each part into Unsafe and not Unsafe. (Placeholder: the figure is not yet drawn.)

under recipe G (sampled). It carries more of Qwen's rise (0.64–0.78).

Answers that open directly also became less safe. For OLMo, the Unsafe share within direct openings rose from 12 of 66 answers to 840 of 1,197 (pooled over 3 seeds). So the first words are not the whole account; compare the shallow-alignment view of Qi et al. (2025). Recipe G keeps more refusal openings. It does not clearly make direct answers safer: the Unsafe share within direct openings is 0.60 under recipe G and 0.68 under recipe B (seed 1101).

This analysis is observational only. We ran no prefill test, and we count positions in words, not tokens.

6. Does content matter? Dark against light fiction

[PART 3: placeholder section. Design only; no data are built and no run has been made.]

Question. Does the topic of benign data change the erosion when the format and length of the data stay the same? Two accounts predict different results. Under a *content* account, dark topics shift what the model is willing to talk about. Under a *format* account, the shape of the data drives the erosion, through the loss on the first answer tokens (§5).

Arms. Three arms on OLMo-2-1B differ in one thing each. Arm D is the Dolly data of §3–4. Arms F-light and F-dark each hold 4,800 rows of human-written fiction from public-domain books (Project Gutenberg, first published in 1930 or earlier), 8 books per arm. Dark books carry the catalogue's horror, Gothic or ghost-story labels. Light books carry its humour, love-story or domestic-fiction labels. The *content test* is F-dark against F-light. The *format test* is F-light against D.

Rows. Each row gives 1 to 3 sentences of context from the book, a fixed instruction ("Continue the story." or a shorter variant), and the original passage as the response. Each passage ends at a sentence boundary near the length of a matched Dolly row, so the three arms share one length profile. Without this match, length could pass for content: short rows put more of a per-row loss on the first answer tokens. The dark arm takes the passages with the highest score on a frozen darkness word list; the light arm takes those with the lowest. No row text is written by a model.

Screens. We drop rows that match a frozen list of real-world procedure patterns, and rows that Granite Guardian flags. The outcome judge never selects training rows. We report both judges' labels on each arm's training rows, as a check that the dark arm is darker. We froze the data rules before any download (Section A).

Judges and reading rule. Qwen3Guard counts depictions of violence as Unsafe, so the HarmBench compliance judge is co-primary here. A content claim needs E and E_{HB} to agree in sign. We will say that content matters if the interval for the difference excludes 0 under both judges. We will say that there is no large content effect if the 90% interval lies inside ± 10 pp. Otherwise the test is inconclusive. With 3 seeds per arm, the minimum detectable difference is roughly 9–10 pp.

Predictions. We will freeze these before any run. P1 (content): E_{HB} is larger for F-dark than for F-light, with the same sign for E. P2 (category): the dark arm's extra Unsafe answers fall mostly in the Violent category. P3 (format): E is smaller for F-light than for D. P4 (mechanism): in every arm, the opening switch carries about half of the change. P5 (judge gap): $E - E_{\mathrm{HB}}$ is larger for F-dark than for F-light.

Recipe and seeds. [COMMON-RECIPE: the main study's common recipe (recipe G at the rung its learning-rate ladder fixes); seeds 1101 to 1103; the sealed untuned baseline reused.]

Results. [CONTENT-RESULTS: Table 3 with E and E-HB per arm and seed; U without Violent-only answers as a check; the reading under the rule above.]

7. What we screened, and limits

Screening log. We list every setting we ran, in time order. All runs use LoRA rank 64 and a recipe like recipe B, unless marked.

1. Qwen2.5-1.5B, Stage 0 recipe (learning rate 2×10^{-4}, 8 rows per step), 3 seeds: mean $E = +81.8$ pp after 3 epochs. It failed the bar that answers must end.

2. Qwen2.5-1.5B, recipe B, seed 1101: $+75.2$ pp. Seed 1102 was lost to a timer and gave no result.

3. Qwen2.5-7B, recipe B, seed 1101: +82.2 pp.
4. Qwen2.5-1.5B, recipe B at rank 16, seed 1101: +67.8 pp.
5. OLMo-2-1B, recipe B, seed 1101: +43.5 pp; then seeds 1102 and 1103.
6. OLMo-2-1B, recipe G, seed 1101: +29.0 pp. We pre-registered the prediction $E_{\mathrm{G}} < E_{\mathrm{B}}$.
7. Llama-3.2-1B, recipe B, seed 1101: +62.5 pp.
8. [QWEN-G, LLAMA-G, G-SEEDS, BT, HAWKINS-LIKE and L1, if run; then the fiction arms.]

Disclosures.

- We fixed the 20–55 pp window for partial erosion after the three Qwen results. No written reason for 20 and 55 exists.
- Seed 1101 chose the setting, and it counts in the OLMo mean. The mean of the fresh seeds (1102 and 1103) is +45.9 pp.
- The OLMo band covers evaluation noise only. The between-seed variance rests on 2 df: the sample SD's 90% range is 0.23–1.73 times the true SD.
- U is a broad policy label. After fine-tuning, the probability of an Unsafe label given no refusal is 0.91–0.94, so U mostly tracks lost refusal.
- Under hard categories only, OLMo's E is about 28 pp. The HarmBench classifier gives $E_{\mathrm{HB}} = +27.8$, $+30.5$ and $+32.8$ pp on seeds 1101–1103 (mean $+30.4$ pp, sampled answers): about two thirds of the Qwen3Guard mean. Its agreement with Qwen3Guard on the fine-tuned runs is $\kappa = 0.65$–0.71.
- Recipe B is not standard practice: a per-row mean loss, a constant learning rate and no warmup. Recipe G is the standard-practice comparison.
- The families differ in chat template, tokeniser (so the same rows give different token counts), post-training pipeline and starting over-refusal (68% for Qwen2.5-1.5B, 17% for OLMo, 11% for Llama and 7% for Qwen2.5-7B).
- The OLMo and Qwen runs used different engine identities. The Llama and OLMo recipe-G runs used a third (Section A).
- The Qwen and Llama results rest on one seed each. The ARC check is uninformative at 1B.
- [HUMAN-READ: a blind human read of about 50 answers, with the rubric fixed first.]

Judge wording. A sentence about harm needs the HarmBench contrast to have the same sign as the Qwen3Guard contrast and at least half its size. If E_{HB} is below 20 pp, no sentence says harm. For OLMo-2-1B under recipe B the rule holds on every seed ($E_{\mathrm{HB}} \geq 20$ pp and at least half of E), so we say harm there. Under recipe G, E_{HB} is below 20 pp on every seed, so we say that recipe G's erosion is lost refusal, not harm, although the same-sign, half-size condition holds. The rule holds for every Qwen and Llama run (E_{HB} from $+35.5$ to $+63.5$ pp, at least 0.61 of E), so we say harm there too.

8. Related work

Benign fine-tuning erodes safety. Qi et al. (2024) showed that fine-tuning aligned models on benign instruction data, such as Alpaca and Dolly, erodes their safety. He et al. (2024) showed that benign rows chosen for their likeness to harmful rows erode safety more than random rows, and that list-format answers raise erosion by themselves. [CITE-CHECK] Kim et al. (2025) show that optimisation choices, such as the learning rate, batch size and number of epochs, drive much of the safety loss.

Family contrasts exist. Hawkins et al. (2026) fine-tune small models from several families with one LoRA recipe ($\alpha/r = 8$, constant learning rate 5×10^{-5}, one epoch, 1,000 synthetic rows seeded from Dolly) over 3 seeds, and measure SORRY-Bench compliance. In English, Gemma-3-1B rises from 30.2% to 65.5%, Llama-3.2-1B from 34.5% to 38.4%, and Qwen3-0.6B falls from 83.2% to 69.8% (their Table 2; they read the Qwen3 drop as a possible collapse). So the same Llama-3.2-1B erodes little in their setup and nears saturation in ours. Their data, judge, starting rate and recipe all differ from ours. Khan et al. (2026) fine-tune four 7–9B families on medical and legal data. They write: "Base model family was likewise an unreliable predictor" (their §4.2.1). Guo et al. (2026) compare Llama-3.1-8B and Qwen2.5-7B, with no OLMo model. On Alpaca, Qwen erodes less (LoRA on 5,000 rows: 9.1% against 44.8% attack success). On Dolly with full fine-tuning, the order flips (73.6% against 66.9%). So at 7–8B, the Qwen-versus-Llama order in Guo et al. (2026) depends on the data. Fraser et al. (2025) fine-tune Llama-3.2-1B on Dolly with 4-bit LoRA for 5 epochs. Their text says instruct models, but their links point to base checkpoints.

Shallow alignment. Qi et al. (2025) argue that safety alignment sits mostly in the first few answer tokens. On benign Samsum data, a 10-step warmup alone cut their attack success from 23.4% to 3.8% (Table 9 of the ICLR version; Table 5 of arXiv v1). They credit their constrained objective mainly for the attack settings. Recipe G includes a warmup, and §5 finds that the opening switch carries about

half of OLMo's erosion.

Content of benign data. Zhao et al. (2025b) fine-tune 95 role-specific LLaMA-3-8B models and find that safety falls further the more villainous the character. Vetter et al. (2026) trace the features behind emergent misalignment to text about villains and harmful agency. Fine-tuning on such human-written documents did not reliably cause it, while synthetic instruction-answer pairs made from the same content did. Narrow fine-tuning can generalise broadly (Betley et al., 2026), and answer format alone can raise erosion (He et al., 2024). Our design in §6 (human-written story continuation) sits between these results.

OLMo-2. We are not aware of prior benign-fine-tuning erosion numbers for OLMo-2 Instruct. Our check found work on refusal geometry in OLMo-2-1B, whose only fine-tuning adds refusals, and work on OLMo-2 7B and 13B with harmful data only [OLMO-PRIOR: cite 2608.25390, 2610.00320]. One web search found no such numbers, and it was not exhaustive.

Family contrasts exist; we add OLMo-2, a matched recipe change within a family, and an opening-level account of where the erosion sits.

9. Conclusion

One recipe can give very different erosion across families at the same size. A standard loss and schedule cut the erosion in one family, so an erosion number must come with its recipe. Report erosion per family and recipe, with the judge's definition and every setting screened. [CONTENT: one sentence on the fiction test, once run.]

Impact Statement

[IMPACT: two to four sentences; the points to cover are in the source.]

References

Betley, J., Tan, D., Warncke, N., Sztyber-Betley, A., Bao, X., Soto, M., Labenz, N., and Evans, O. Emergent misalignment: Narrow finetuning can produce broadly misaligned LLMs, 2026. URL https://arxiv.org/abs/2502.17424.

Clark, P., Cowhey, I., Etzioni, O., Khot, T., Sabharwal, A., Schoenick, C., and Tafjord, O. Think you have solved question answering? Try ARC, the AI2 reasoning challenge. arXiv preprint arXiv:1803.05457, 2018. URL https://arxiv.org/abs/1803.05457.

Conover, M., Hayes, M., Mathur, A., Xie, J., Wan, J., Shah, S., Ghodsi, A., Wendell, P., Zaharia, M., and Xin, R. Free Dolly: Introducing the world's first truly open instruction-tuned LLM. Databricks blog, 2023. URL https://www.databricks.com/blog/2023/04/12/dolly-first-open-commercially-viable-instruction-tuned-llm.

Cui, J., Chiang, W.-L., Stoica, I., and Hsieh, C.-J. OR-Bench: An over-refusal benchmark for large language models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 11515–11542. PMLR, 2025. URL https://proceedings.mlr.press/v267/cui25a.html.

Fraser, K. C., Dawkins, H., Nejadgholi, I., and Kiritchenko, S. Fine-tuning lowers safety and disrupts evaluation consistency, 2025. URL https://arxiv.org/abs/2506.17209.

Grattafiori, A., Dubey, A., Jauhri, A., et al. The Llama 3 herd of models. arXiv preprint arXiv:2407.21783, 2024. URL https://arxiv.org/abs/2407.21783.

Guo, Y., Chen, X., Zhang, S., Wang, X., and Tang, H. When safety routing breaks: Understanding alignment fragility under benign fine-tuning. arXiv preprint arXiv:2609.01455, 2026. URL https://arxiv.org/abs/2609.01455. Accepted to the Findings of EMNLP 2026.

Hawkins, W., Rawal, K., Rystrøm, J., Tsirtsis, S., Fu, Z., Warren, G., Brown, R., Delaney, E., Wachter, S., Mittelstadt, B., and Russell, C. The heterogeneous safety impacts of benign multilingual fine-tuning, 2026. URL https://arxiv.org/abs/2606.28843.

He, L., Xia, M., and Henderson, P. What is in your safe data? Identifying benign data that breaks safety. In *First Conference on Language Modeling*, 2024. URL https://openreview.net/forum?id=Hi8jKh4HE9.

Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., and Chen, W. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022. URL https://openreview.net/forum?id=nZeVKeeFYf9.

Khan, E. B., Winecoff, A., Bogen, M., and Hadfield-Menell, D. Safety drift after fine-tuning: Evidence from high-stakes domains, 2026. URL https://arxiv.org/abs/2604.24902.

Kim, M., Kwak, J. M., Alssum, L., Ghanem, B., Torr, P., Krueger, D., Barez, F., and Bibi, A. Rethinking safety in LLM fine-tuning: An optimization perspective. In *Second Conference on Language Modeling*, 2025. URL https://openreview.net/forum?id=ZnOoEA2nDn.

Mazeika, M., Phan, L., Yin, X., Zou, A., Wang, Z., Mu, N., Sakhaee, E., Li, N., Basart, S., Li, B., Forsyth, D., and Hendrycks, D. HarmBench: A standardized evaluation framework for automated red teaming and robust refusal. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 35181–35224. PMLR, 2024. URL https://proceedings.mlr.press/v235/mazeika24a.html.

Padhi, I., Nagireddy, M., Cornacchia, G., et al. Granite Guardian. arXiv preprint arXiv:2412.07724, 2024. URL https://arxiv.org/abs/2412.07724.

Qi, X., Zeng, Y., Xie, T., Chen, P.-Y., Jia, R., Mittal, P., and Henderson, P. Fine-tuning aligned language models compromises safety, even when users do not intend to! In *The Twelfth International Conference on Learning Representations*, 2024. URL https://openreview.net/forum?id=hTEGyKf0dZ.

Qi, X., Panda, A., Lyu, K., Ma, X., Roy, S., Beirami, A., Mittal, P., and Henderson, P. Safety alignment should be made more than just a few tokens deep. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=6Mxhg9PtDE.

Qwen, Yang, A., Yang, B., et al. Qwen2.5 technical report. arXiv preprint arXiv:2412.15115, 2024. URL https://arxiv.org/abs/2412.15115.

Röttger, P., Kirk, H., Vidgen, B., Attanasio, G., Bianchi, F., and Hovy, D. XSTest: A test suite for identifying exaggerated safety behaviours in large language models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 5377–5400. Association for Computational Linguistics, 2024. URL https://aclanthology.org/2024.naacl-long.301/.

Team OLMo, Walsh, P., Soldaini, L., Groeneveld, D., et al. 2 OLMo 2 furious. arXiv preprint arXiv:2501.00656, 2024. URL https://arxiv.org/abs/2501.00656.

Vetter, C., Kaczér, D., Flek, L., and Mai, F. Data attribution of emergent misalignment with persona features, 2026. URL https://arxiv.org/abs/2608.11025.

Zhao, H., Yuan, C., Huang, F., et al. Qwen3Guard technical report. arXiv preprint arXiv:2510.14276, 2025a. URL https://arxiv.org/abs/2510.14276.

Zhao, W., Hu, Y., Deng, Y., Guo, J., Sui, X., Han, X., Zhang, A., Zhao, Y., Qin, B., Chua, T.-S., and Liu, T. Beware of your Po! Measuring and mitigating AI safety risks in role-play fine-tuning of LLMs, 2025b. URL https://arxiv.org/abs/2502.20968.

A. Hashes

Table 3. Model, data, recipe and engine hashes, and the configuration SHA-256s. Highlighted cells are pending.

Item	Revision or SHA-256
OLMo-2-0425-1B-Instruct, revision	[HASH]
Qwen2.5-1.5B-Instruct, revision	`989aa798...` [full hash]
Qwen2.5-7B-Instruct, revision	`a09a3545...` [full hash]
Llama-3.2-1B-Instruct, revision	[HASH]
Prepared Dolly data, one manifest per family	[HASHES]
Recipe configurations (B, G, Stage 0, bt, Hawkins-like)	[HASHES]
Engine identity: OLMo, recipe B	`55253874...` [full hash]
Engine identity: Qwen2.5	`c52c0ebc...` [full hash]
Engine identity: Llama-3.2, and OLMo recipe G	`e49a978c...` [full hash]
Logit-scale prediction for recipe G on Qwen (§4)	[PREDICTION-RECORD: commit and time]
Fiction data rules v1, frozen 2026-10-08 20:32 UTC, before any download (§6)	`17fe8f2495dd86a64969a5e59accf6ad6ed32ae1f096e5f2c75e22cf0a8bb84b`

B. Per-seed tables

[SEEDS: per-seed U, E, SE and band for OLMo recipe B (seeds 1101 to 1103) and recipe G; Granite per run; capped answers and screen counts.]

C. Judge tables

[JUDGES: Qwen3Guard, Granite and HarmBench per run; kappa between the judges; the hard-category cut; Unsafe labels on the over-refusal bank (5% untuned, 22 to 29% after fine-tuning).]

D. Run log and costs

[RUNS: every run in time order, with the pod, GPU type, hours and cost in USD; failed and lost runs included.]

E. Reproducibility

[REPRO: scripts and their hashes; the frozen opening word list (lexicon bbbc71e9, full hash); the bars files; the configurations.]